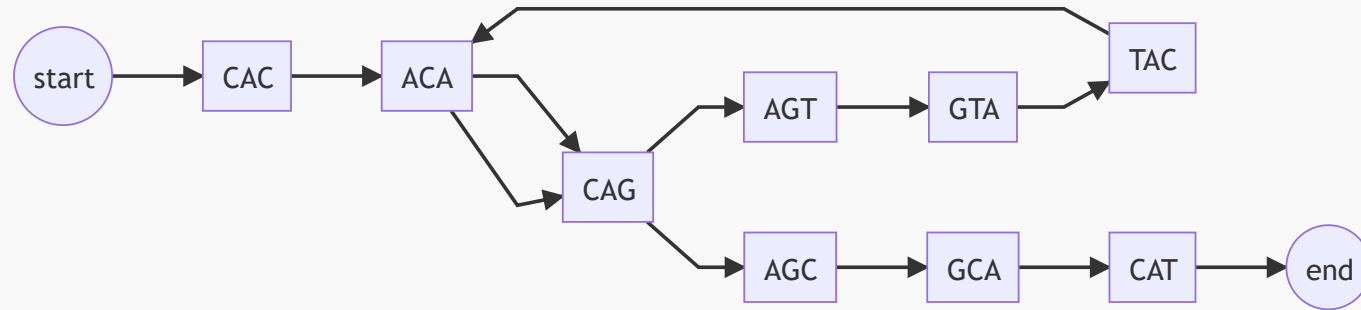


## BIOL8706: Dividing and conquering sequence alignment using De Bruijn Graphs



- Student: Richard Morris
- Huttley lab, Australian National University
- Supervisors: Gavin Huttley, Vijini Mallawaarachchi



# Introduction to sequence alignment

Given we can sequence genomes of different organisms.

Sequence A: **ATGCATAC** Sequence B: **ATGTAC**

We can compare sequences. But first we have to align these sequences to identify common regions

<b>A</b>	<b>T</b>	<b>G</b>	<b>C</b>	<b>A</b>	<b>T</b>	<b>A</b>	<b>C</b>
↕	↕	↕	<b>X</b>	<b>X</b>	↕	↕	↕
<b>A</b>	<b>T</b>	<b>G</b>	<b>—</b>	<b>—</b>	<b>T</b>	<b>A</b>	<b>C</b>


To investigate this difference, we need to identify regions that are different, and regions that are similar. To do that we will put these two sequences in a data structure called a partial order graph

## Sequence as a partial order graph


Our alignment

<b>A</b>	<b>T</b>	<b>G</b>	<b>C</b>	<b>A</b>	<b>T</b>	<b>A</b>	<b>C</b>
↕	↕	↕	<b>X</b>	<b>X</b>	↕	↕	↕
<b>A</b>	<b>T</b>	<b>G</b>	<b>—</b>	<b>—</b>	<b>T</b>	<b>A</b>	<b>C</b>


can be represented as the following partial order graph, showing each node and the direction of the alignment.



# Extracting regions from the partial order graph



By collecting together adjacent nodes with the same number of edges we can simplify that to



**Now we can make some claims about which regions are present in both sequences**


eg: If those regions encoded for genes, then we can make some claims about organism genotype.

	ATG	CA	TAC
Sequence A	✓	✓	✓
Sequence B	✓	∅	✓

# Why is multiple sequence alignment (MSA) important?

Alignment of eg: a viral genome allows us to:

- Identify conserved regions for vaccine/drug development
- Identify changes in function to make predictions about the virus' behaviour
- Identify and prepare for emerging variants



Alignment of S mutation points of SARS-CoV-2 variants

# Why is MSA so computationally expensive?

- An exact solution has an order complexity of  $O(L^n)$ 
    - **L** is the length of the sequence
    - **n** is the number of sequences
-

# MSA for SARS-CoV-2 genomes?

## SARS-CoV-2

- length: **~29,903** bp
- number: **over 5 million** (as of March 2022) <sup>1</sup>
- $O(29,903^{\text{over 5 million}})$  is a **very large number**

**Required: a method to align large numbers of small sequences**




Fig 1: Artists rendition of SARS-CoV-2


<sup>1</sup> [doi.org/10.1038/s41588-022-01033-y](https://doi.org/10.1038/s41588-022-01033-y) | Fig 1 [doi.org/10.7875/togopic.2020.199](https://doi.org/10.7875/togopic.2020.199)

# MSA for great apes genomes?

## The great apes

- length: **~3 billion bp**
- number: 5
- $O(3\text{Billion}^5)$  is also a very large number.
- However great ape genomes are 97+% identical<sup>1</sup>

**Required: a method to identify the few different regions in very long similar sequences**



The family tree of great apes

<sup>1</sup> [citation needed](#)



## Computation requires energy

### The human genome moonshot alignment

- Calculations performed by chatGPT
- length: ~3 billion bp
- number: ~8 billion
- $O(3B^{8B})$  calculations
- Supercomputers perform  $10^{18}$  FLOPs per second
- at an efficiency of ~30 billion FLOPs/watt
- assuming the order unit is one FLOP

=  $3,000,000,000^{8,000,000,000} / 30,000,000,000$  watts

The milky way over it's 13.6 Billion year lifetime it can be calculated has only generated  $3.8 \times 10^{37}$  watts

**So without a more efficient algorithm we'll need the energy of 216 million more galaxies to perform this calculation**



## **Required: a more efficient method to align**

- large numbers of small sequences
- small numbers of very similar long sequences

# Sequence alignment order complexity

## Pairwise sequence alignment

- Compare every letter in one sequence to every letter in the other
- order complexity of  $O(mn)$ 
  - where **m** and **n** are lengths of the sequences

## Multiple sequence alignment (MSA)

- Perform a pairwise alignment of every sequence to every other sequence
- order complexity of  $O(L^n)$ 
  - where **L** is the length of the sequences
  - **n** is the number of sequences


## Pairwise sequence alignment methods: $O(mn)$

- Needleman-Wunsch algorithm: global alignment for highly similar sequences
  - scoring system that penalises gaps and mismatches
- Smith-Waterman algorithm: better for local alignment to find conserved domains
  - allows for alignment to reset when the score falls to 0

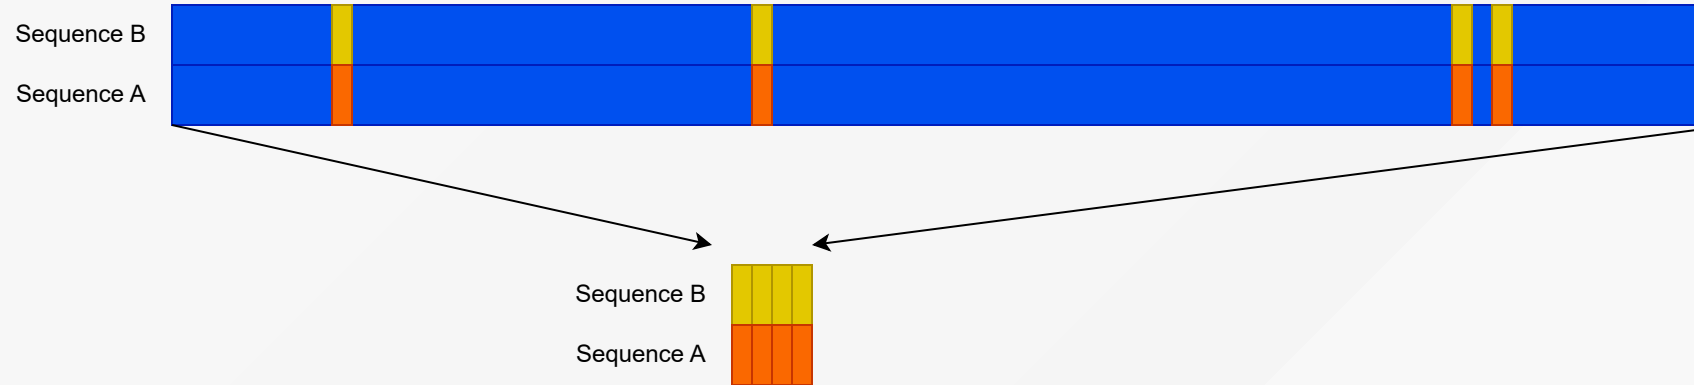
## Multiple sequence alignment (MAS) strategies

- Exact alignment
  - $O(L^n)$
- Progressive alignment eg: ClustalW
  - create a guide tree  $O(n^2.L^2)$
  - Progressively align pairs most closely related to profiles, and then align profiles  $O(n^2.L)$
- Iterative methods eg: MUSCLE, T-Coffee, MAAFT
  - $O(n^2.L^2)$
  - create an preliminary fast less accurate alignment
  - iteratively improve alignment using some scoring function
  - Complete when some convergence criterion is met
- Hidden markov models  $O(nL) + O(LM)$  (M is the number of states in the model)
  - eg: HMMER
  - create a statical model of the transition between states
  - Determine likely alignment based on the model

## The chicken and egg problem of MSA



## What if we could quickly remove regions that are similar?



**We'd be able to focus our computational resources on just the regions that are different.**

# Sequence alignment using De Bruijn Graphs

This work builds on the work by Xingjian Leng in a 12 month undergraduate research project in 2022, under the supervision of Dr. Yu Lin and Prof. Gavin Huttley.

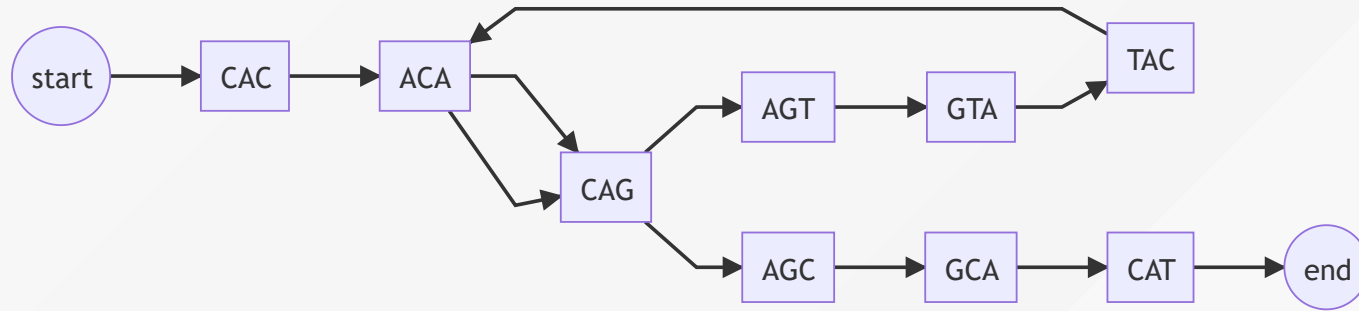
That project focused on the alignment of closely related viral genomes, with a particular emphasis on SARS-CoV-2. The method is based on the construction and utilization of de Bruijn graphs for both pairwise and multiple sequence alignment tasks.



## De Bruijn graphs


A De Bruijn graph is a directed graph that represents unique overlapping subsequences (or k-mers) at the nodes. This structure is an efficient way to identify sequence overlaps, and common regions.

Building a De Bruijn graph has an order complexity of  $O(nL)$



## Overlapping k-mers

Consider the DNA sequence *CACAGTACGGCAT* when broken into 3 character overlapping subsequences (or 3-mers) looks like this:



# De Bruijn graphs


When we represent that as a de Bruijn graph it looks like this:



## A second sequence

Consider we want to align that sequence  $CACAGTAC\boxed{G}GCAT$  to the very similar sequence  $CACAGTAC\boxed{T}CGCAT$


Which as a De Bruijn graph looks like this:



## De Bruijn pairwise alignment




**Sequence A:**




**Sequence B:**

If we combine both sequences into a single de Bruijn graph, we can easily identify the regions that are similar and the regions that are different.



## Resolving the graph

We can collect nodes with 2 edge, or 1 edge into single nodes, and we can see the regions that are similar and the regions that are different.




Now we can use a traditional algorithm to align the regions  $AC\boxed{G}GC$  and  $AC\boxed{T}GC$ , and we've reduced  $O(14^2)$  down to  $O(5^2)$  = **7.8x** less work.

## De Bruijn multiple sequence alignment

And we can extend this to multiple sequences. Consider aligning the following sequences

CACAGTACGGCAT CACAGTACTGCAT CACAGTACTGGAGCAT & CACAGTACTGATGCAT



Now we've reduced  $O(13 \times 13 \times 16 \times 16)$  down to  $O(6 \times 6 \times 8 \times 8) = \mathbf{18.8x}$  less work

## Reducing the horizontal complexity of the problem

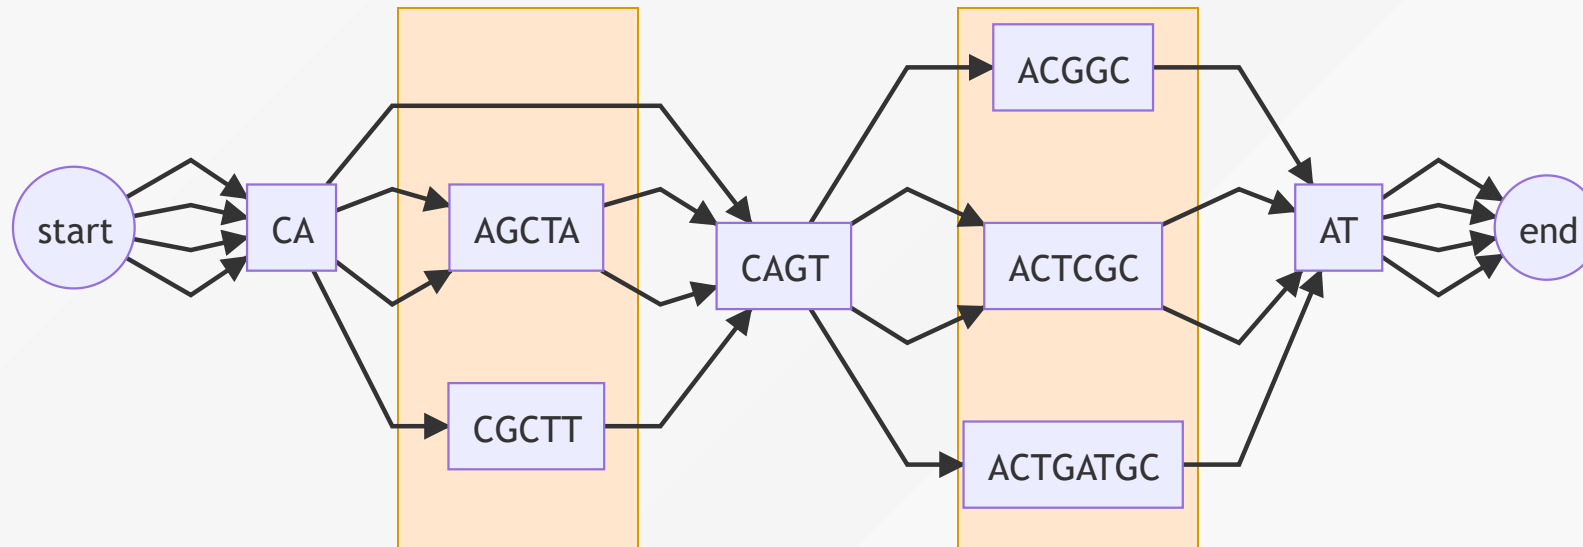
- Horizontal component of the problem is the length (L) of the sequences to be aligned
- recall an exact alignment has an order complexity of  $O(L^n)$
- if we reduce the length of the sequences we need to align we reduce L

**How about n?**



## Reducing the vertical complexity of the problem

- Vertical component of the problem is the number of sequences to be aligned ( $n$ )



- Any matches or deletions reduces the number of sequences we need to align

## Project aims

- Investigate the use of De Bruijn graphs to identify regions of dissimilarity for traditional alignment algorithms
- Build a python library for implementing De Bruijn Graph Multi-sequence alignment
  - Resolve the De Bruijn graph to a partial order graph to reduce horizontal complexity
  - Add bitmaps to nodes identifying sequences that contain that node
  - Use bitmaps to reduce vertical complexity
  - investigate using bitmaps for identifying reverse compliment regions
  - investigate using bitmaps for identifying sequence distance
- Investigate measures of performance and accuracy
- Stretch goals
  - Quantify **performance** against traditional methods
  - Quantify **accuracy** against traditional methods

# Results

**TBD...**


# Discussion

**TBD...**

## Future directions

Investigate the potential of using De Bruijn Graphs to;

- identify reverse compliment regions from a DBG
- identify phylogeny from a DBG using bitmaps to measure distance
- investigate phylogeny of conserved regions in species with lateral gene flow



# Thanks

- Gavin Huttley
- Yu Lin
- Vijini Mallawaarachchi
- Xinjian Leng
- Huttley lab

# Questions