BIOL8706: Dividing and conquering sequence alignment using De Bruijn Graphs

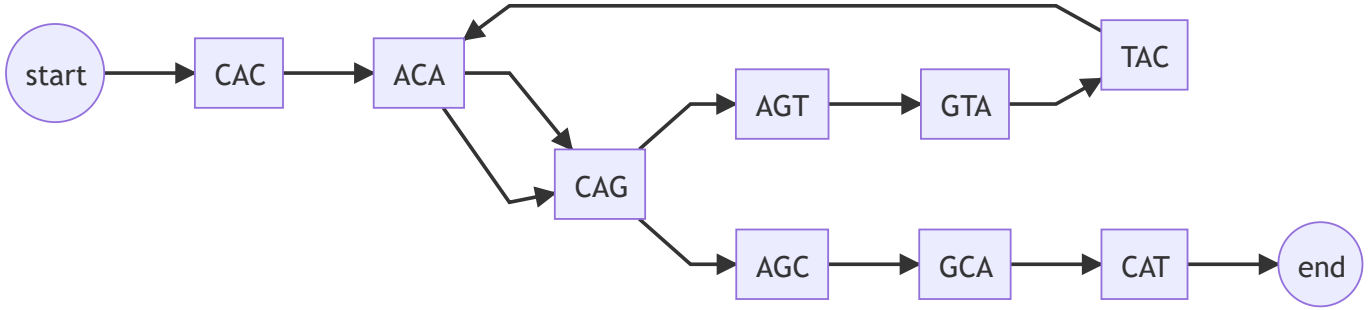

- Student: Richard Morris
- Huttley lab, Australian National University
- Supervisors: Gavin Huttley, Vijini Mallawaarachchi

Australian National University

Why should I care about sequence alignment?

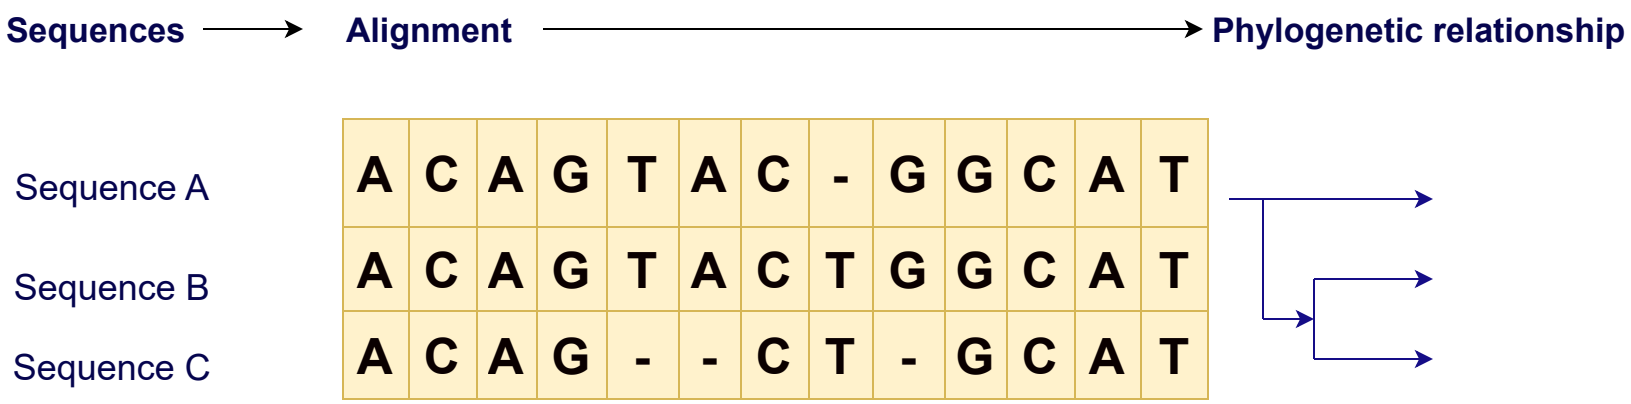

- Sequence alignment arranges sequences of DNA, RNA or Protein to identify regions of similarity
- Uncover evolutionary relationships between sequences

Consider Rubisco

- one of the most abundant proteins on Earth

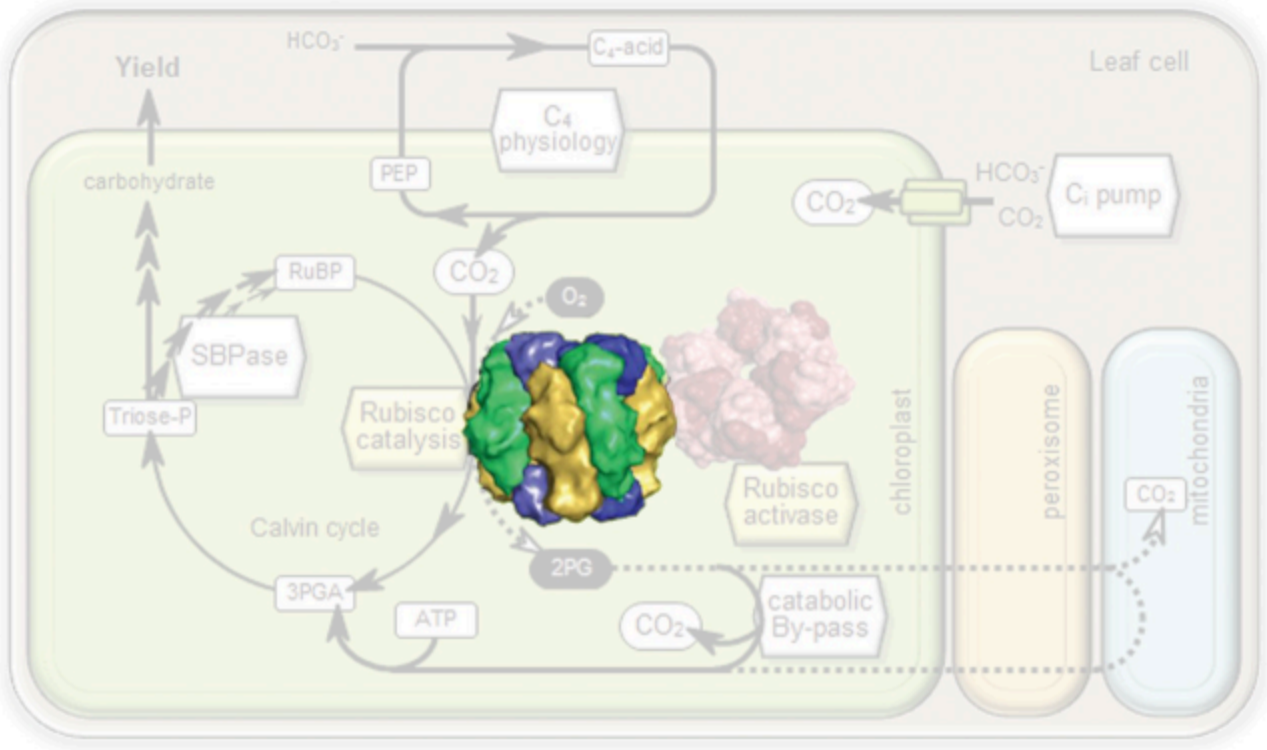

Parry et al 2012 doi.org/10.1093/jxb/ers336

Consider Rubisco

- one of the most abundant proteins on Earth
- **essential component of photosynthesis**

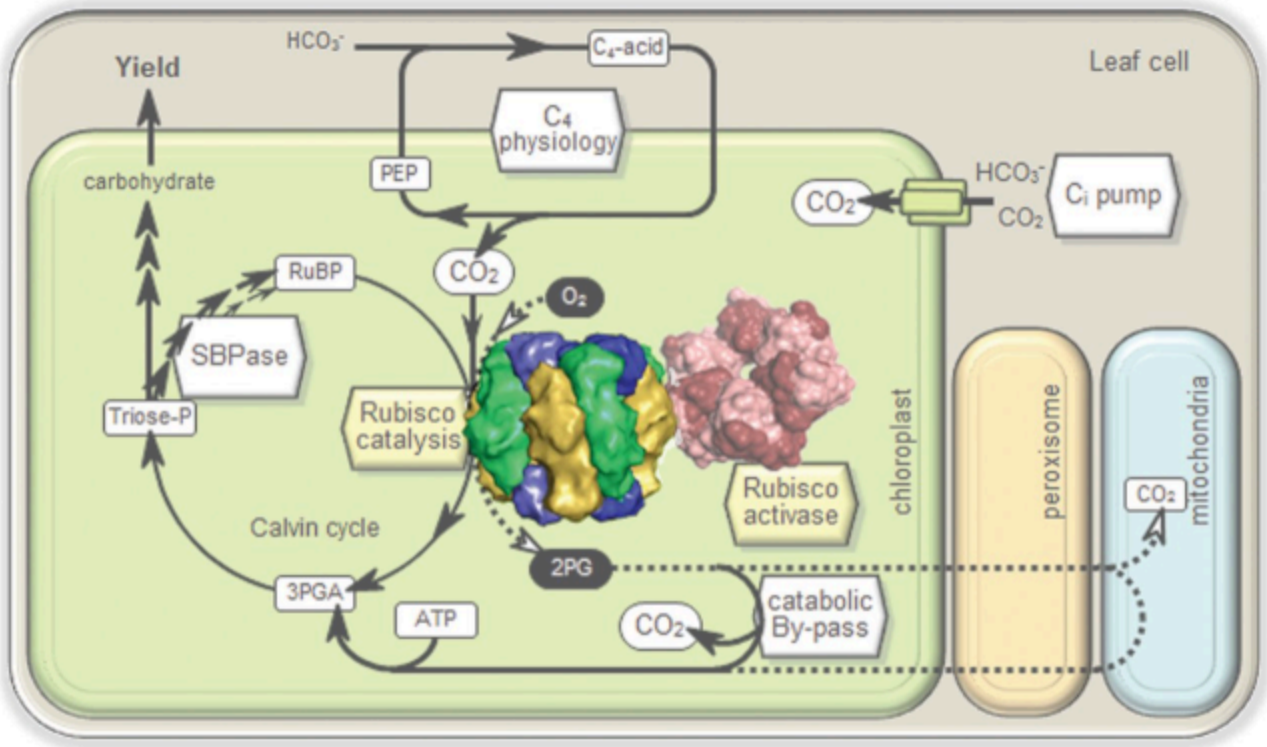

Parry et al 2012 doi.org/10.1093/jxb/ers336

Consider Rubisco

- one of the most abundant proteins on Earth
- essential component of photosynthesis
- **converts CO2 to organic carbon**

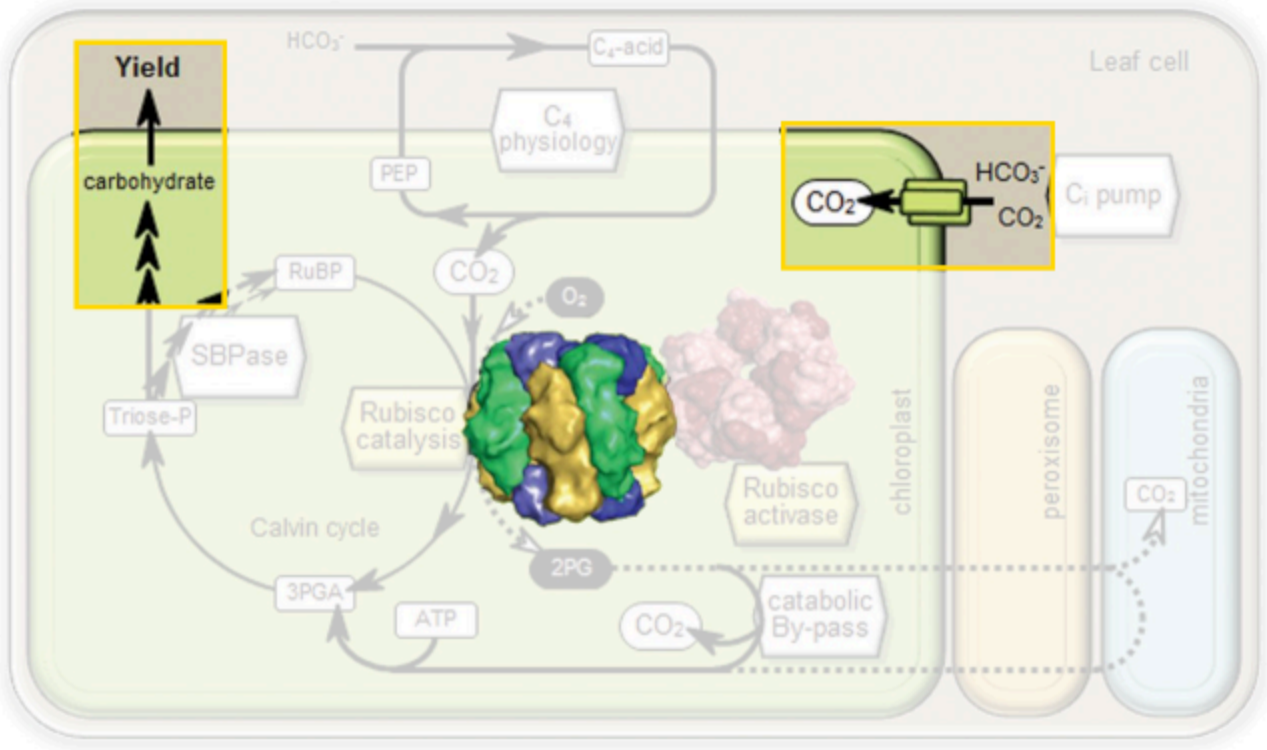

Parry et al 2012 doi.org/10.1093/jxb/ers336

Evolution of Rubisco

- Genomic sequencing has identified Rubisco-like proteins in 3 kingdoms of life
- Phylogenetic analysis supports the existence of 3 clades of Rubisco

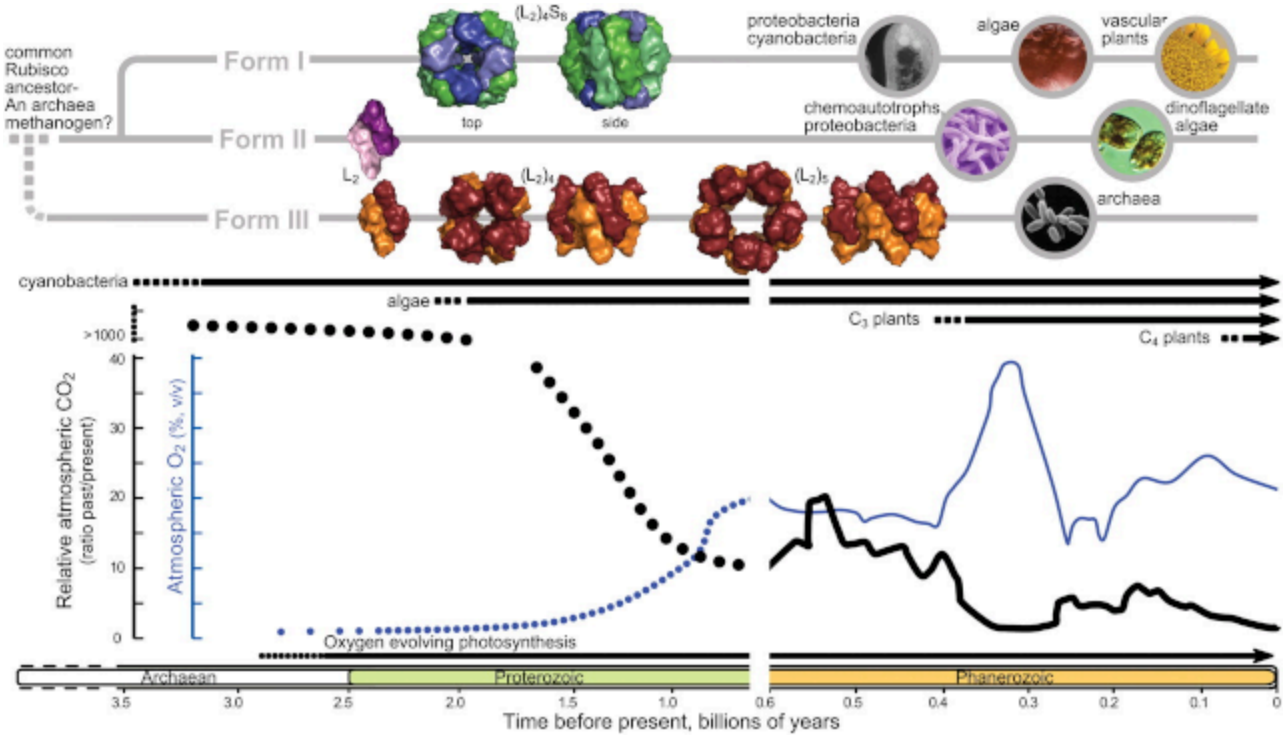

Whitney et al 2012 doi.org/10.1104/pp.110.164814

How we build a phylogeny from extant sequences

Sequences → **Alignment** → **Phylogenetic relationship**

Sequence A	A	C	A	G	T	A	C	-	G	G	C	A	T
Sequence B	A	C	A	G	T	A	C	T	G	G	C	A	T
Sequence C	A	C	A	G	-	-	C	T	-	G	C	A	T

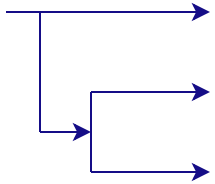

- Sequence alignment is the first step in building a phylogeny
- Exhaustive alignment compares every character in each sequence with every character in every other sequence

Exhaustive alignment is costly

> “ Code slows as data grows ”
>
> *Ned Batchelder*

What is sequence alignment

Alignment of eg: a viral genome allows us to:

- Identify conserved regions for vaccine/drug development
- Identify changes in function to make predictions about the virus' behaviour
- Identify and prepare for emerging variants

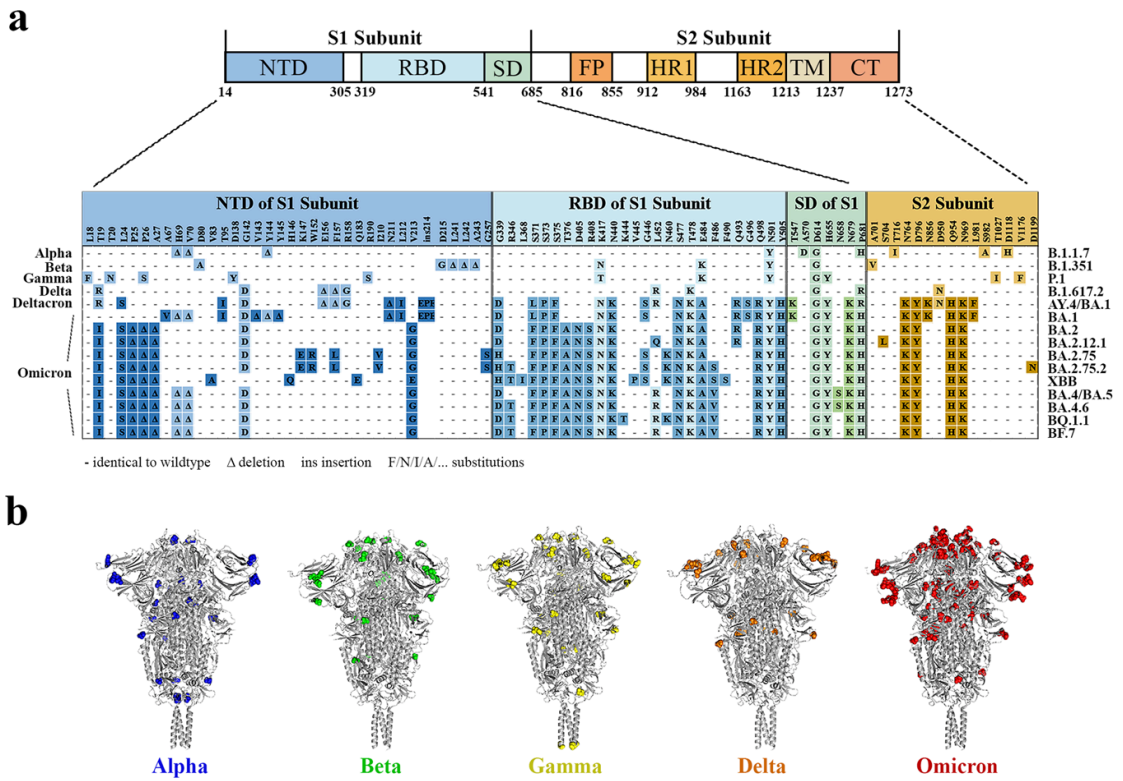

Alignment of S mutation points of SARS-CoV-2 variants

Why is MSA so computationally expensive?

- An exhaustive solution has an order complexity of $O(L^n)$
  - **L** is the length of the sequence
  - **n** is the number of sequences

MSA for SARS-CoV-2 genomes?

SARS-CoV-2

- length: **~29,903** bp
- number: **over 5 million** (as of March 2022) [1]
- $O(29,903^{\text{over 5 million}})$ is a **very large number**

Required: a method to align large numbers of small sequences

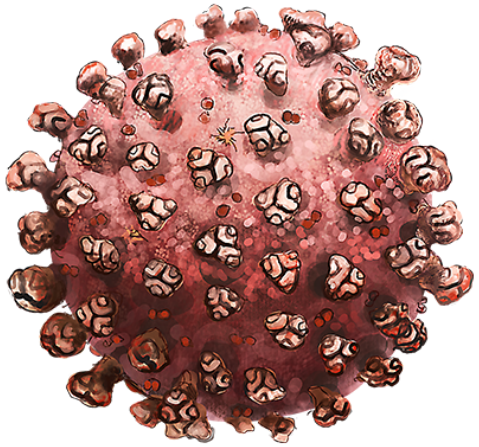

Fig 1: Artists rendition of SARS-CoV-2

[1] doi.org/10.1038/s41588-022-01033-y | Fig 1 doi.org/10.7875/togopic.2020.199

MSA for great apes genomes?

The great apes

- length: **~3 billion** bp
- number: 5
- $O(3Billion^5)$ is also a very large number.
- However great ape genomes are 97+% identical[1]

Required: a method to identify the few different regions in very long similar sequences

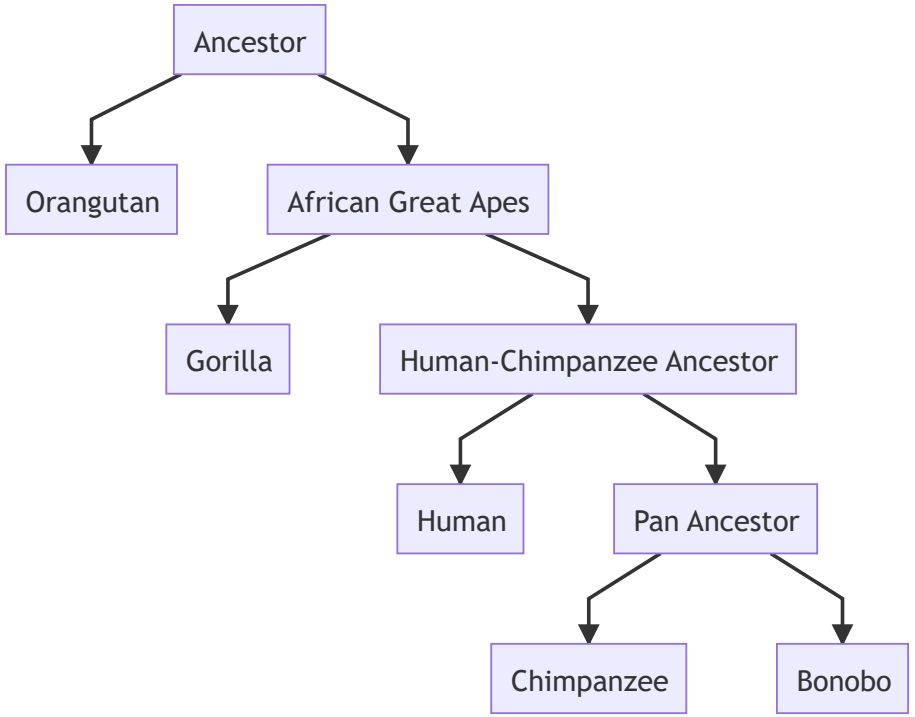

The family tree of great apes

[1] citation needed

Alignment takes energy

- Sequence alignment requires **computation**
- Computation requires **energy**

Required: a more efficient method to align

- large numbers of small sequences
- small numbers of very similar long sequences

Sequence alignment order complexity

Pairwise sequence alignment

- Compare every letter in one sequence to every letter in the other
- order complexity of $O(mn)$
  - where **m** and **n** are lengths of the sequences

Multiple sequence alignment (MSA)

- Perform a pairwise alignment of every sequence to every other sequence
- order complexity of $O(L^n)$
  - where **L** is the length of the sequences
  - **n** is the number of sequences

- scoring system that penalises gaps and mismatches

- Smith-Waterman algorithm: better for local alignment to find conserved domains
  - allows for alignment to reset when the score falls to 0

-1 mismatch, -2 gap (δ)

Where $F(i,j) = \max$ of the following

$$\nwarrow F(i-1,j-1) + s(A_i, B_j), \quad \text{(match/mismatch)}$$
$$\Uparrow F(i-1,j) + \delta, \quad \text{(deletion)}$$
$$\Leftarrow F(i,j-1) + \delta, \quad \text{(insertion)}$$

	gap	A	G	C	A	
gap	***0***	⇐-2	⇐-4	⇐-6	⇐-8	⇐
A	⇑-2	⇖ ***1***	⇐ -1	⇐-3	⇐-5	⇐
C	⇑-4	⇑-1	⇖0	⇖ ***0***	⇐-2	⇐
G	⇑-6	⇑-3	⇖0	⇖-1	⇖ ***1***	⇐
A	⇑-8	⇖⇑-5	⇑-2	⇖-1	⇑-1	⇖
A	⇑-10	⇖⇑-7	⇑-4	⇖⇑-3	⇖-2	⇖

backtrace from bottom right selecting the

Multiple sequence alignment (MAS) strategies

Dividing and conquering sequence alignment using De Bruijn Graphs

- Pairwise alignment of each possible pair
  - $\binom{n}{2} \times O(L^2) = \frac{n(n-1)}{2} \times O(L^2) = O(n^2.L^2)$
- Progressive alignment eg: ClustalW
  - create a guide tree
  - Progressively align pairs most closely related to profiles, and then align profiles
- Iterative methods eg: MUSCLE, T-Coffee, MAAFT
  - create an preliminary fast less accurate alignment
  - iteratively improve alignment using some scoring function
  - Complete when some convergence criterion is met
- Hidden markov models $O(nL) + O(LM)$ (M is the number of states in the model)
  - eg: HMMER
  - create a statical model of the transition between states

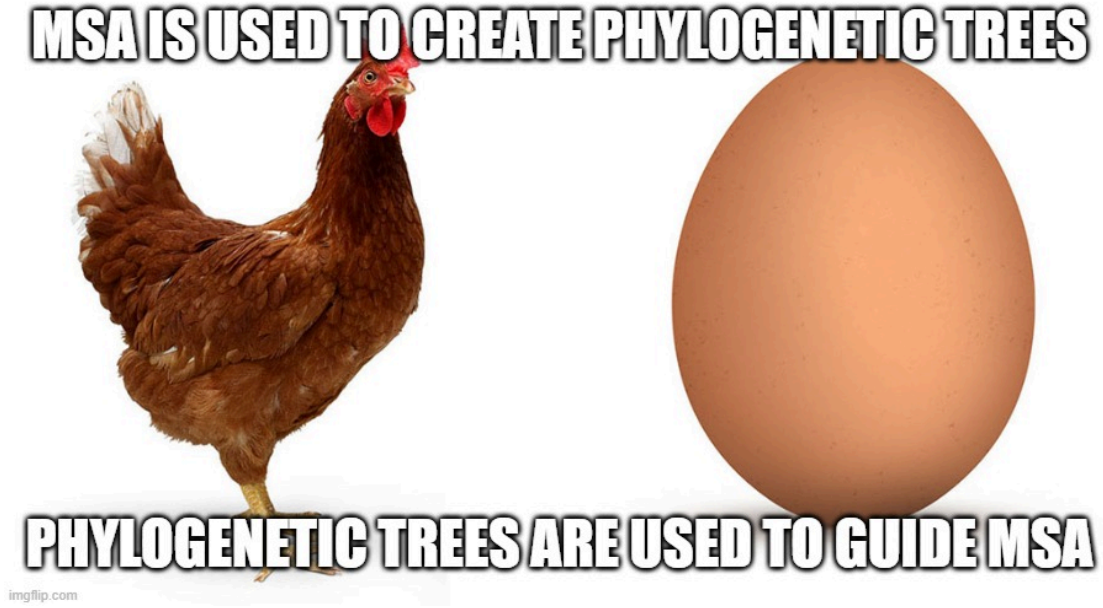

- **Alignment needs trees**

To align sequences accurately and efficiently, one should know the phylogenetic relationships among the sequences to better guide the alignment process

- **Trees need alignment**

to infer a phylogenetic tree accurately, one

What if we could quickly remove regions that are similar?

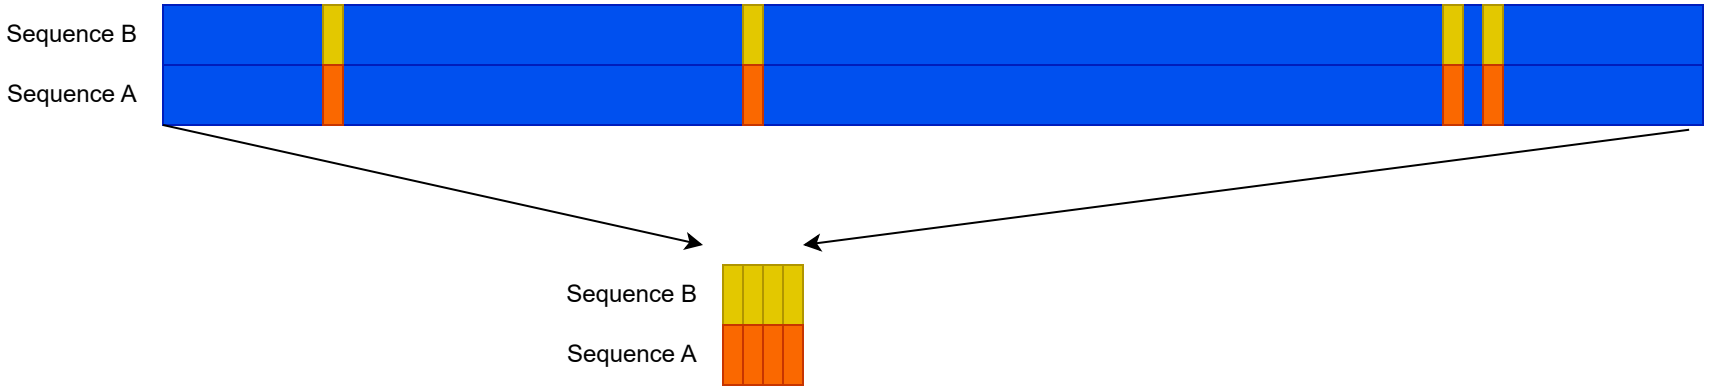

We'd be able to focus our computational resources on just the regions that are different.

Sequence alignment using De Bruijn Graphs

This work builds on the work by Xingjian Leng in a 12 month undergraduate research project in 2022, under the supervision of Dr. Yu Lin and Prof. Gavin Huttley.

That project focused on the alignment of closely related viral genomes, with a particular emphasis on SARS-CoV-2. The method is based on the construction and utilization of de Bruijn graphs for both pairwise and multiple sequence alignment tasks.

De Bruijn graphs

A De Bruijn graph is a directed graph that represents unique overlapping subsequences (or k-mers) at the nodes. This structure is an efficient way to identify sequence overlaps, and common regions.

Building a De Bruijn graph has an order complexity of $O(nL)$

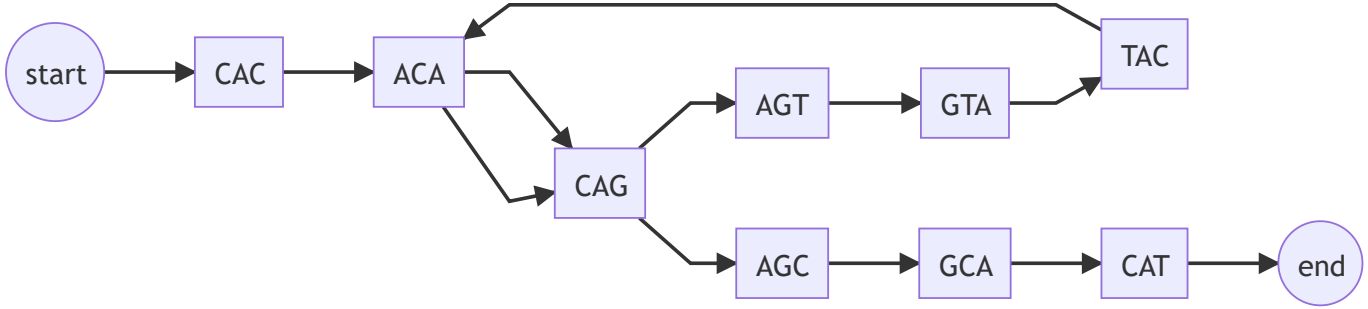

Overlapping k-mers

Consider the DNA sequence $CACAGTACGGCAT$ when broken into 3 character overlapping subsequences (or 3-mers) looks like this:

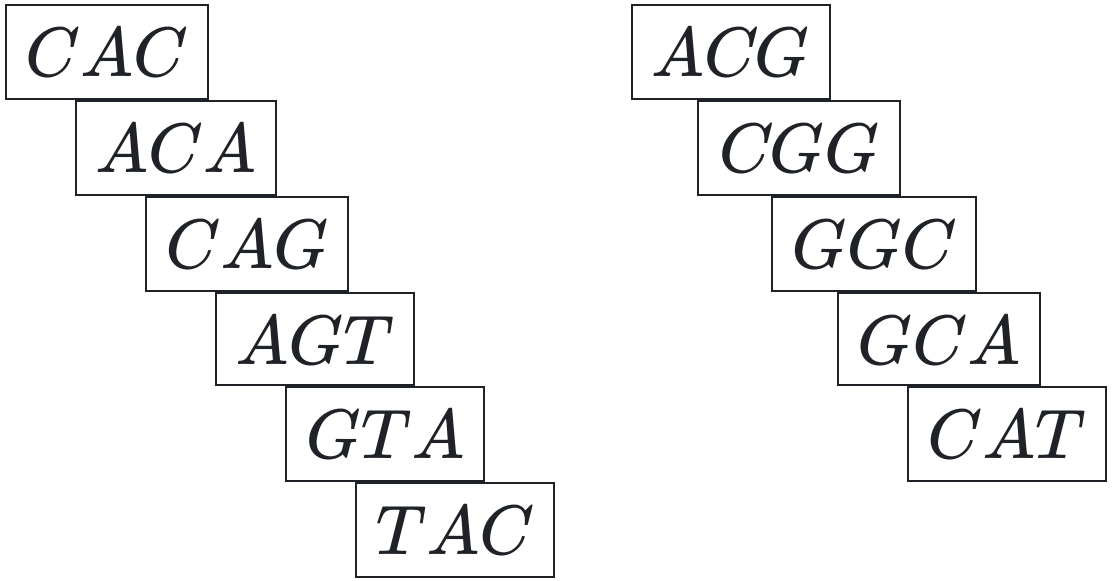

De Bruijn graphs

When we represent that as a de Bruijn graph it looks like this:

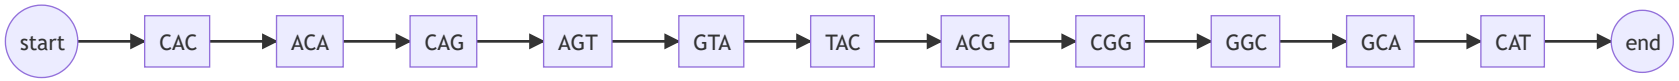

A second sequence

Consider we want to align that sequence $CACAGTAC\boxed{G}GCAT$ to the very similar sequence $CACAGTAC\boxed{T}CGCAT$

Which as a De Bruijn graph looks like this:

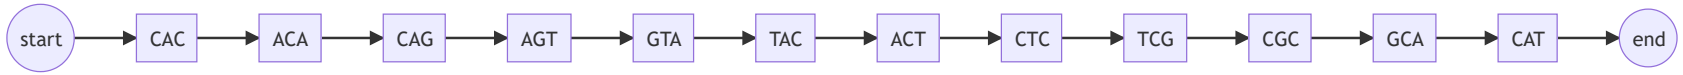

De Bruijn pairwise alignment

Sequence A:

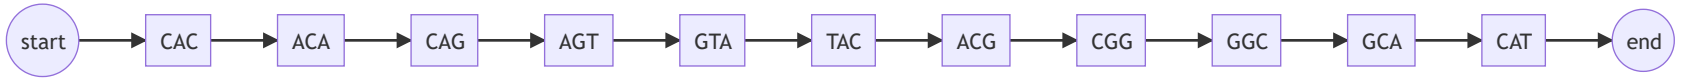

Sequence B:

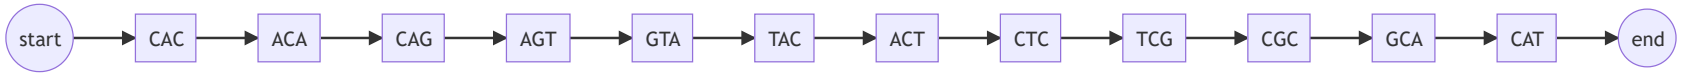

If we combine both sequences into a single de Bruijn graph, we can easily identify the regions that are similar and the regions that are different.

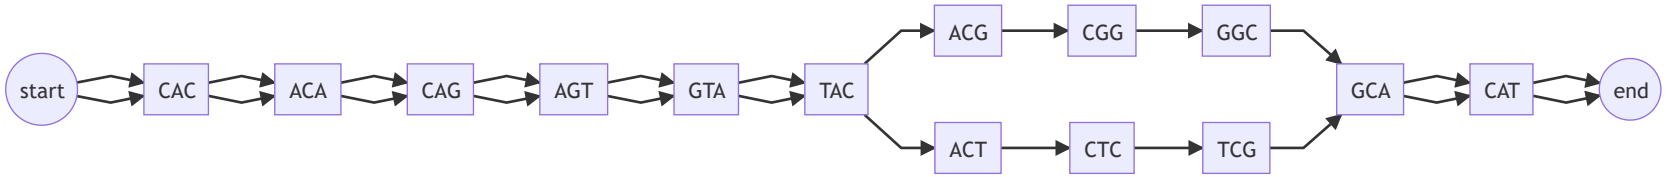

Resolving the graph

We can collect nodes with 2 edge, or 1 edge into single nodes, and we can see the regions that are similar and the regions that are different.

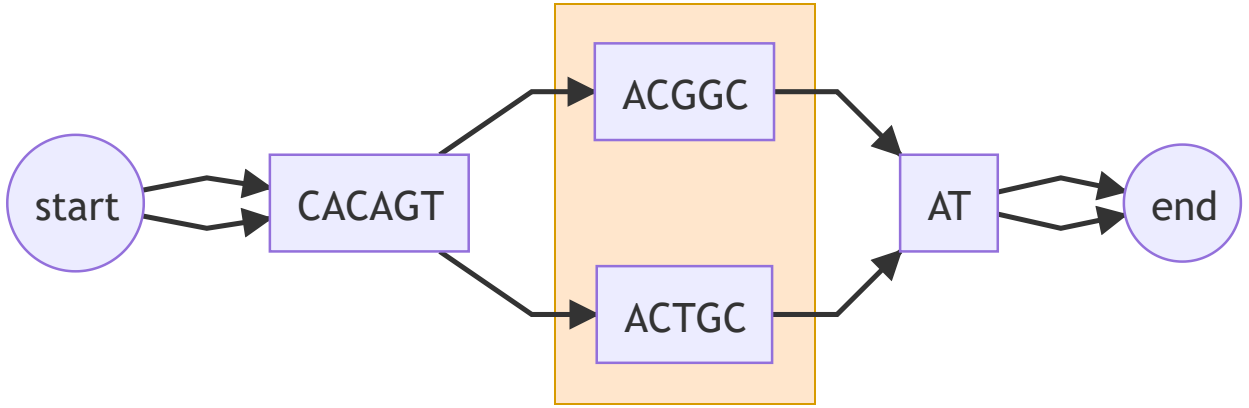

Now we can use a traditional algorithm to align the regions $AC\boxed{G}GC$ and $AC\boxed{T}GC$, and we've reduced $O(14^2)$ down to $O(5^2)$ = **7.8x** less work.

De Bruijn multiple sequence alignment

And we can extend this to multiple sequences. Consider aligning the following sequences

`CACAGTACGGCAT` `CACAGTACTGCAT` `CACAGTACTGGAGCAT` & `CACAGTACTGATGCAT`

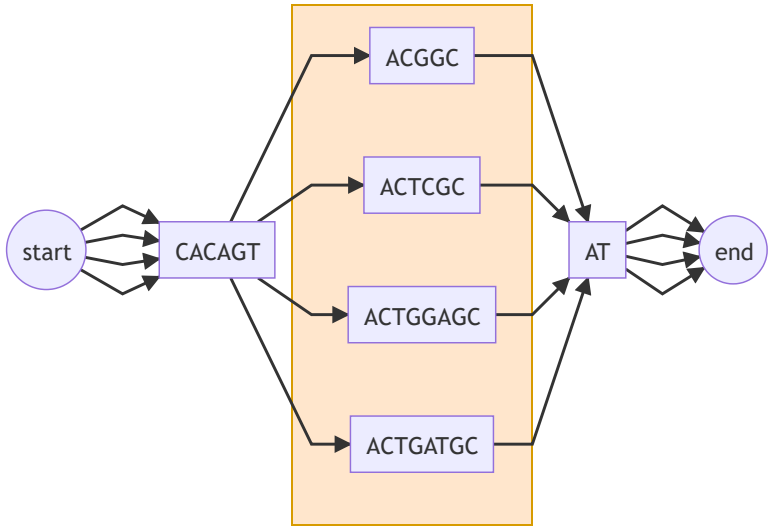

Now we've reduced O(13x13x16x16) down to O(6x6x8x8) = **18.8x** less work

Reducing the horizontal complexity of the problem

- Horizontal component of the problem is the length (L) of the sequences to be aligned
- recall an exact alignment has an order complexity of $O(L^n)$
- if we reduce the length of the sequences we need to align we reduce L

How about n?

Reducing the vertical complexity of the problem

- Vertical component of the problem is the number of sequences to be aligned (n)

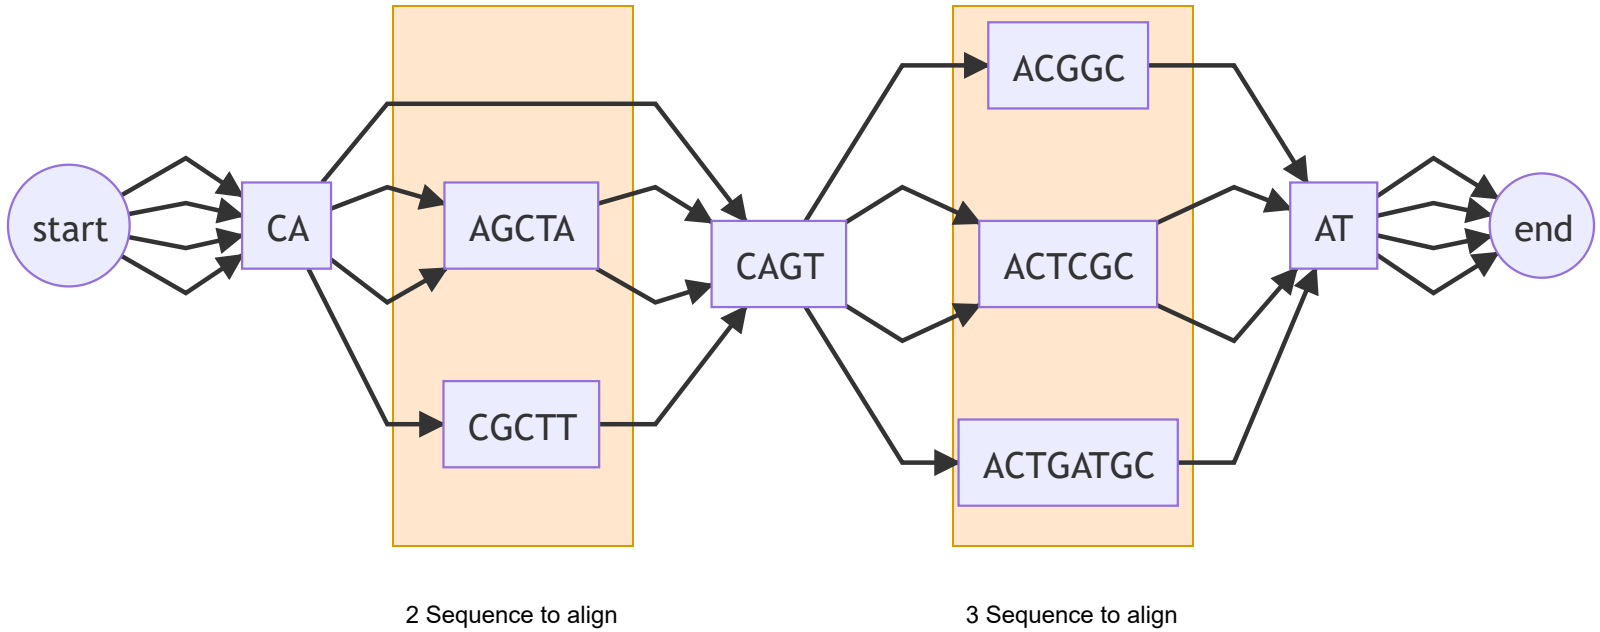

- Any matches or deletions reduces the number of sequences we need to align

    - Resolve the De Bruijn graph to a partial order graph to reduce horizontal complexity
    - Convert matched pairs of bubble edges to profiles to reduce vertical complexity
    - Develop unit tests that verify the correctness of the library against edge case sequence alignments
        - long sequences
        - numerous sequences
        - cyclic sequences
        - bubbles within bubbles
        - sequential bubbles
- Develop statistics for any set of sequences, for a range of possible kmer lengths
    - Time complexity
        - how many alignment operations are required for exact, progressive, 1 dimensional de Bruin graphs, and 2 dimensional de Bruin graphs
    - Memory use ratio

Results: projected memory use ratio

Memory complexity is the amount of actual memory required to stor a sequence divided by the amount of memory required for the original sequences

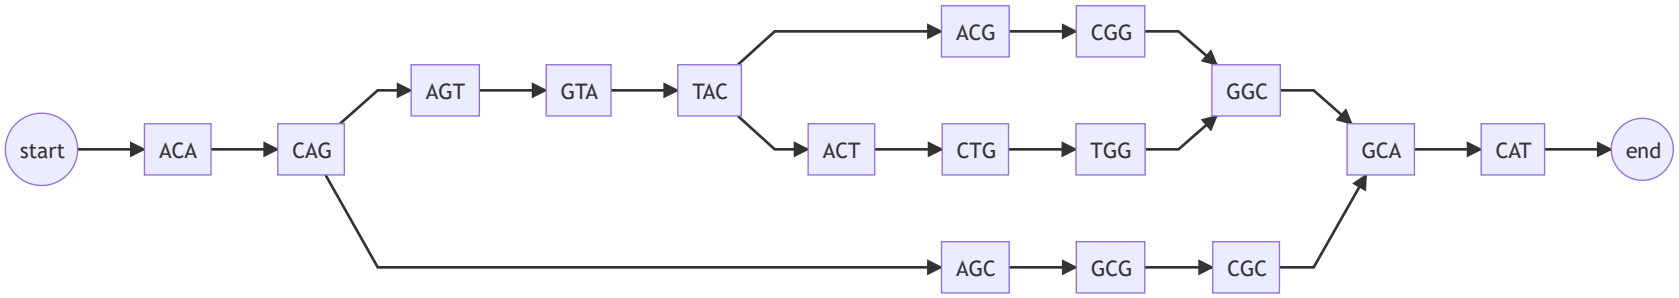

Converted into a partial order graph

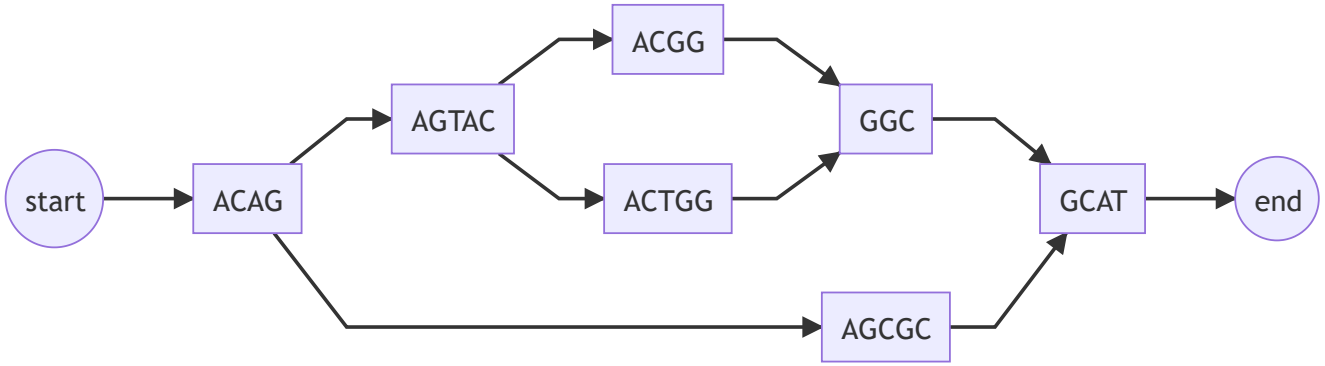

- Memory complexity = $30/34$

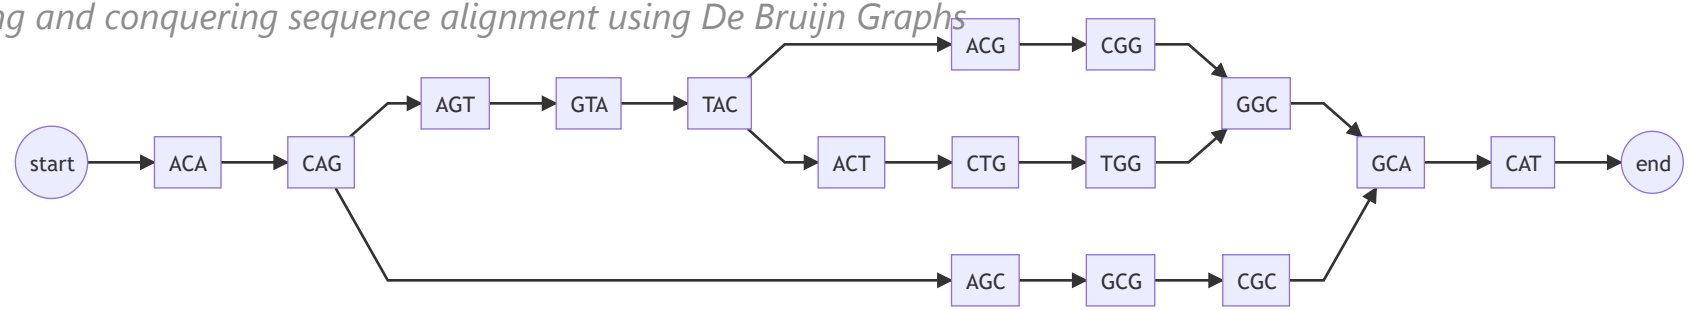

Converted into a partial order graph

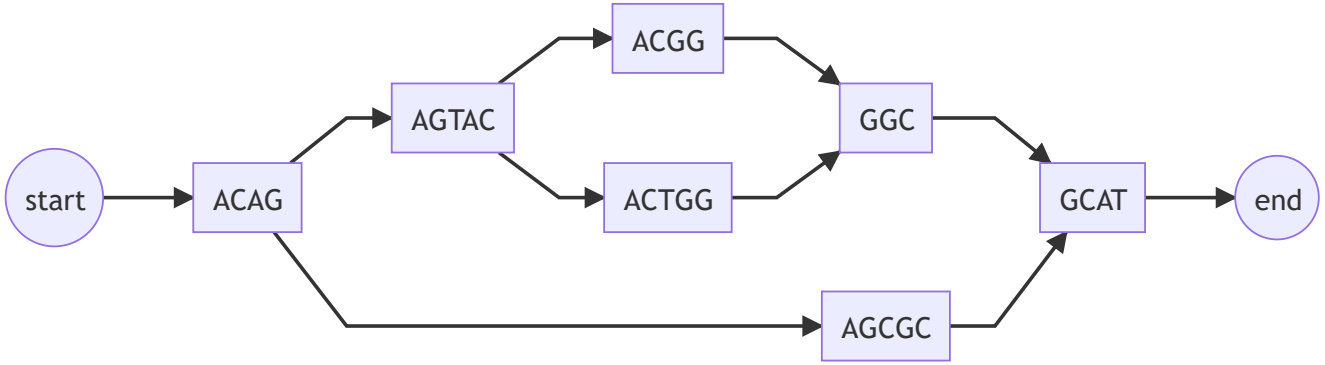

- Exact alignment = $13 \times 12 \times 9 = 1404$
- Progressive alignment = $13 \times 12 + 12 \times 9 + 13 \times 9 = 381$
- 1 dimensional de Brujin Graph simplification then progressive alignment = $9 \times 8 + 8 \times 5 + 9 \times 5 = 157$

Results: Projected compression ratio

Compression ratio is the average length of the payload of partial order graph (POG) nodes divided by the kmer length used to construct the de Brujin graph the POG was derived from.

A higher compression ratio indicates that the kmer size chosen to construct the de Brujin graph from those sequences captured more information about common subsequences.

Given the POG derived from a de Brujin graph with k=3.

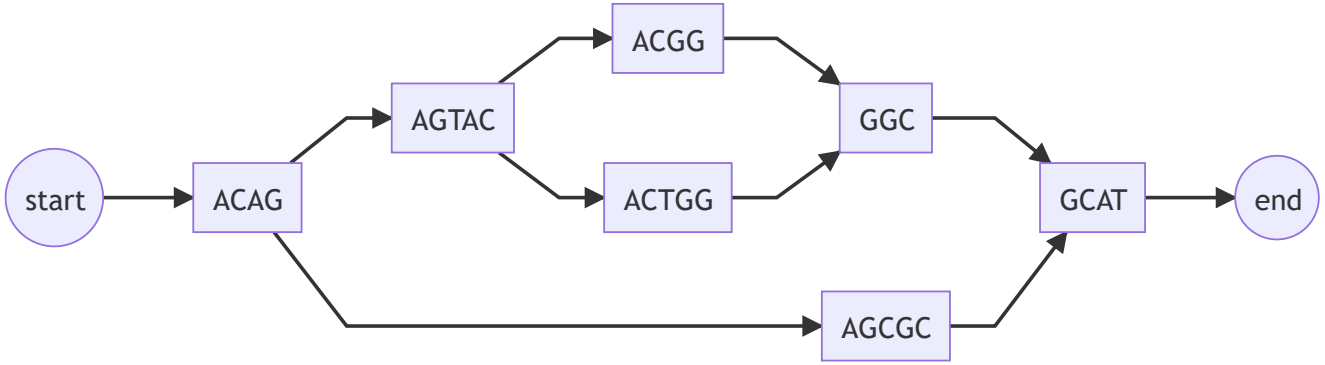

The compression ratio is the mean payload 4.2857 over the original kmer length of 3 $= 1.4286$

phage				

kmer = 6

Genomes	Exact	Progressive	1D dBG	2D dB
BRCA1 56 species				
BRCA1 primates				
SARS-CoV-2				
IBD phage				
Tara				

Results: Calculated compression ratio from alignable sequences

- BRCA1 genes in 56 species (citation needed)
- BRCA1 genes in primates (citation needed)
- SARS-CoV-2 genomes (citation needed)
- IBD phage components (https://doi.org/10.1016/j.cell.2015.01.002)
- Tara oceans phage components (https://doi.org/10.1126/science.1261605)

cyclic sequences

Dividing and conquering sequence alignment using De Bruijn Graphs

```
def test_pog_cycle(output_dir: Path):
    dbg = dbg_align.DeBrujinGraph(3,cogent3.DNA)
    dbg.add_sequence({
        "seq1": "ACAGTACGGCAT",
        "seq2": "ACAGTACTGGCAT",
        "seq3":"ACAGCGCGCAT" # contains cycle
        })
    with open(output_dir / "cycle.md", "w") as f:
        f.write("```mermaid\n")
        f.write(dbg.to_mermaid())
        f.write("```")
    assert dbg.has_cycles()
    assert len(dbg) == 3
    assert dbg.names() == ["seq1", "seq2", "seq3"]
    assert dbg["seq1"] == "ACAGTACGGCAT"
    assert dbg["seq2"] == "ACAGTACTGGCAT"
    assert dbg["seq3"] == "ACAGCGCGCAT" # contains cycle

    dbg.to_pog()
    # write mermaid out to testout folder
    with open(output_dir / "cycle_compressed.md", "w") as f:
        f.write("```mermaid\n")
        f.write(dbg.to_mermaid())
        f.write("```")
```

Discussion

Future directions

Investigate the potential of using de Bruijn Graphs to;

- identify reverse compliment regions from a dBG
- identify genetic distance and infer phylogeny from a dBG

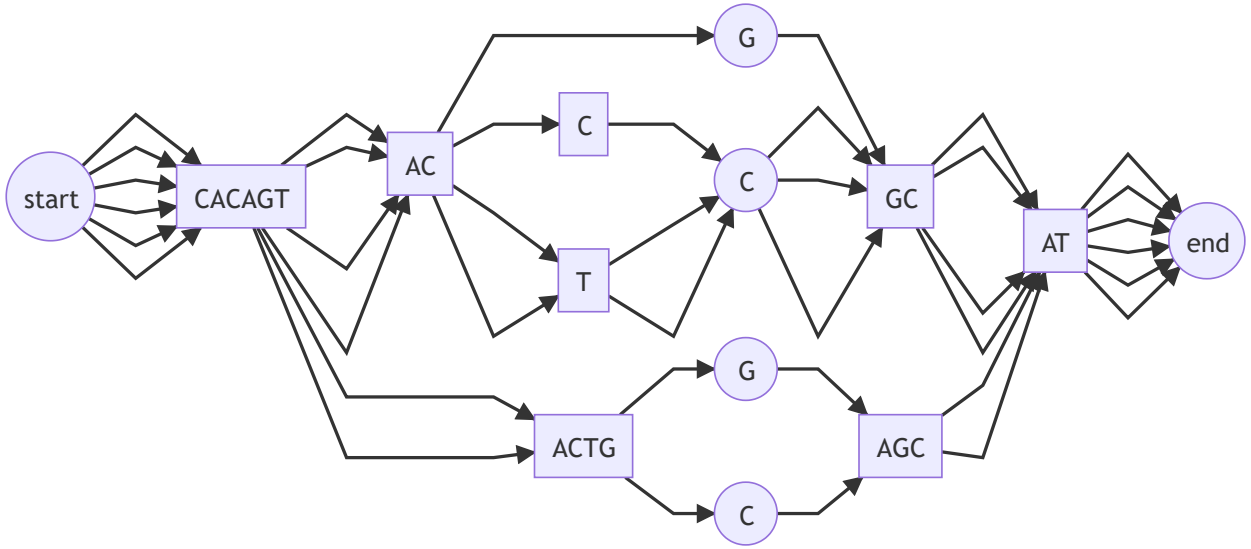

Thanks

- Gavin Huttley
- Yu Lin
- Vijini Mallawaarachchi
- Xinjian Leng
- Huttley lab

Questions