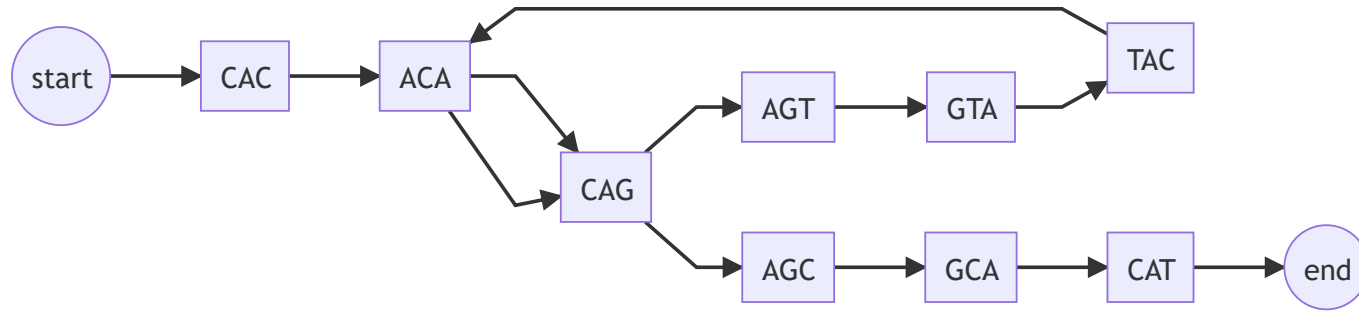


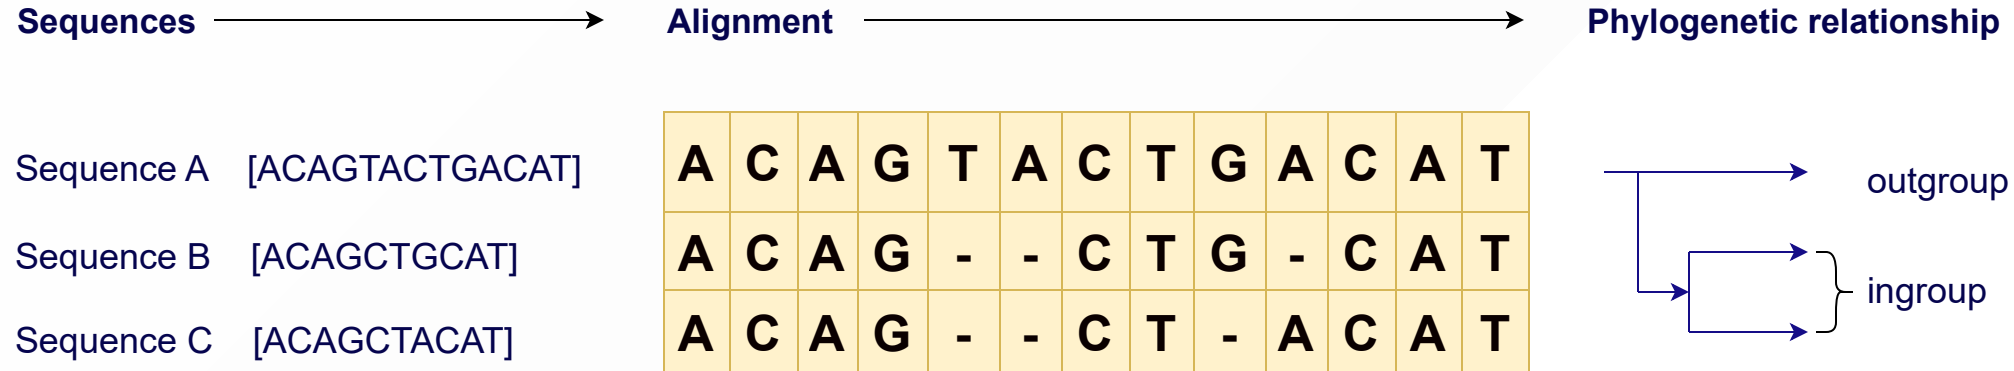
# BIOL8706: Dividing and conquering sequence alignment using De Bruijn Graphs



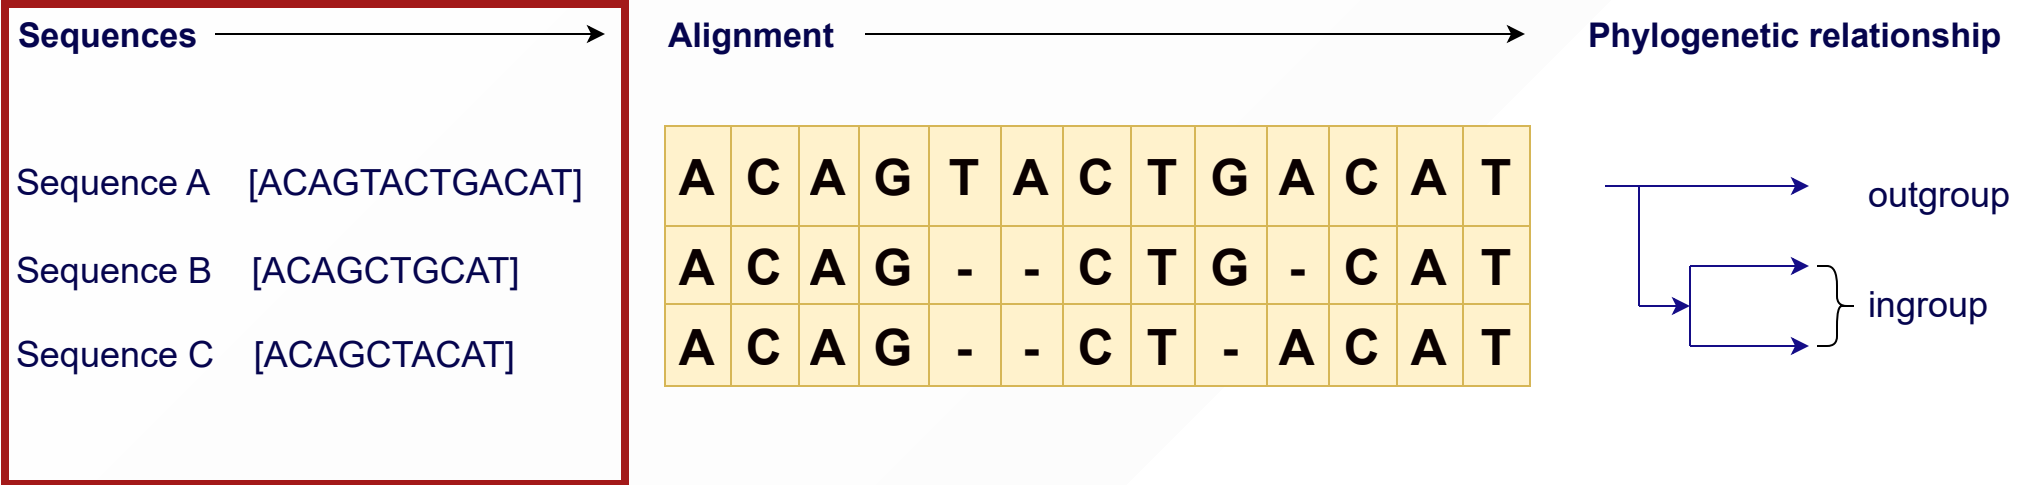
- Student: Richard Morris
- Huttley lab, Australian National University
- Supervisors: Gavin Huttley, Vijini Mallawaarachchi



# Sequence alignment



# Sequence alignment



We start with a set of DNA sequences (Note all different lengths)

# Sequence alignment

Sequences →

Sequence A [ACAGTACTGACAT]


Sequence B [ACAGCTGCAT]

Sequence C [ACAGCTACAT]

Alignment →

A	C	A	G	T	A	C	T	G	A	C	A	T
A	C	A	G	-	-	C	T	G	-	C	A	T
A	C	A	G	-	-	C	T	-	A	C	A	T

Phylogenetic relationship



We align those sequences

# Sequence alignment

Sequences

Sequence A [ACAGTACTGACAT]


Sequence B [ACAGCTGCAT]

Sequence C [ACAGCTACAT]

Alignment

A	C	A	G	T	A	C	T	G	A	C	A	T
A	C	A	G	-	-	C	T	G	-	C	A	T
A	C	A	G	-	-	C	T	-	A	C	A	T

Phylogenetic relationship



By lining up regions that are similar

# Sequence alignment

Sequences →

Sequence A [ACAGTACTGACAT]


Sequence B [ACAGCTGCAT]

Sequence C [ACAGCTACAT]

Alignment →

A	C	A	G	T	A	C	T	G	A	C	A	T
A	C	A	G	-	-	C	T	G	-	C	A	T
A	C	A	G	-	-	C	T	-	A	C	A	T

Phylogenetic relationship



Noting those that are different

# Sequence alignment

Sequences → Alignment →


Sequence A [ACAGTACTGACAT]

Sequence B [ACAGCTGCAT]

Sequence C [ACAGCTACAT]

A	C	A	G	T	A	C	T	G	A	C	A	T
A	C	A	G	-	-	C	T	G	-	C	A	T
A	C	A	G	-	-	C	T	-	A	C	A	T

Phylogenetic relationship



And we can infer evolutionary relationships between those sequences

- ingroup (1 letter different)
- outgroup (3 letters different)

# Sequence alignment


Sequences → Alignment →

Sequence A [ACAGTACTGACAT]

Sequence B [ACAGCTGCAT]

Sequence C [ACAGCTACAT]

A	C	A	G	T	A	C	T	G	A	C	A	T
A	C	A	G	-	-	C	T	G	-	C	A	T
A	C	A	G	-	-	C	T	-	A	C	A	T



**And we can infer evolutionary relationships between those sequences**

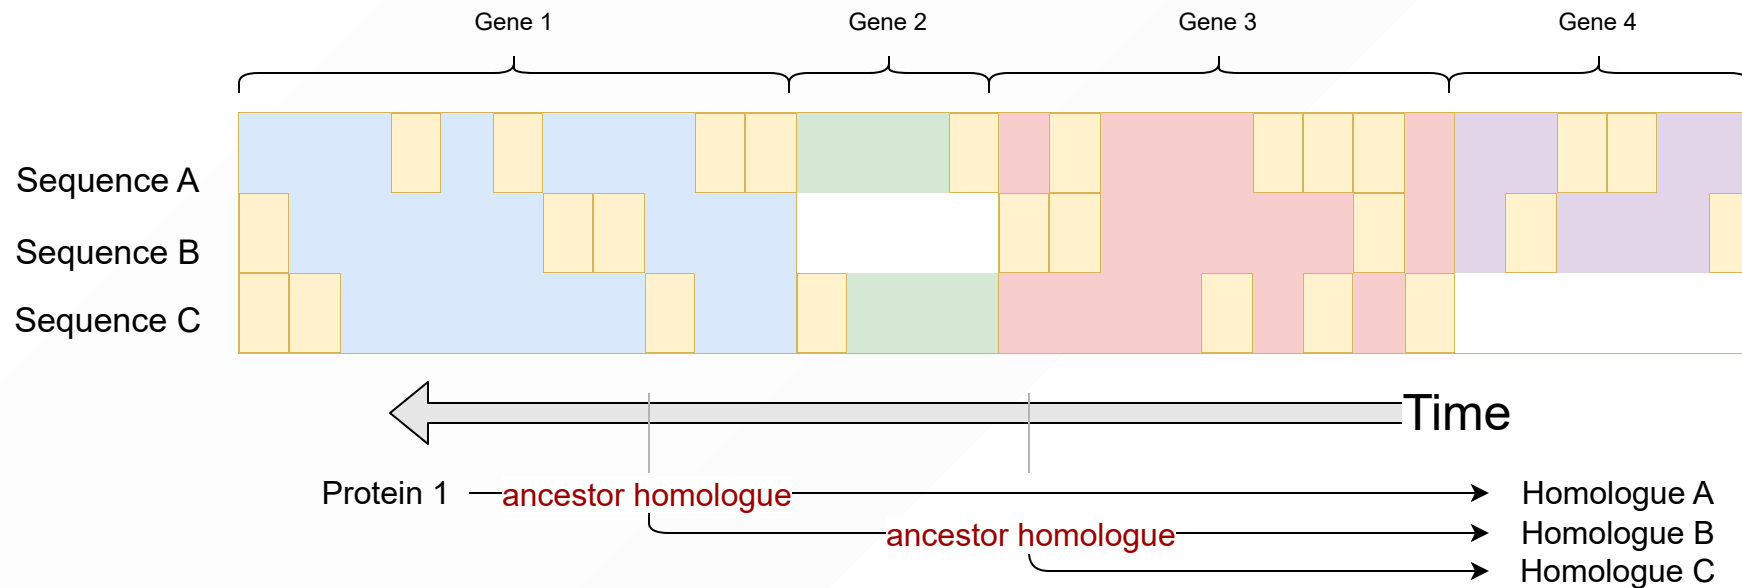
- ingroup (1 letter different)
- outgroup (3 letters different)
- extinct common ancestor sequence
- how long ago sequences diverged



# The central dogma of biology

DNA  $\xrightarrow{\text{transcribe}}$  RNA  $\xrightarrow{\text{translate}}$  Protein

If we can align homologous genes, we can infer homologous proteins




Sequence alignment is a **time machine** for homologous proteins

## **3 big ideas that require sequence alignment**

- **design crops for better yield**
- **predict the trajectory of a virus**
- **understand our own evolution**


# Consider RuBisCO

- The most abundant protein on Earth
- essential component of photosynthesis
- primary role is to convert  $\text{CO}_2$  to organic carbon




# Consider RuBisCO

- The most abundant protein on Earth
- essential component of photosynthesis
- primary role is to convert CO<sub>2</sub> to organic carbon



# Consider RuBisCO

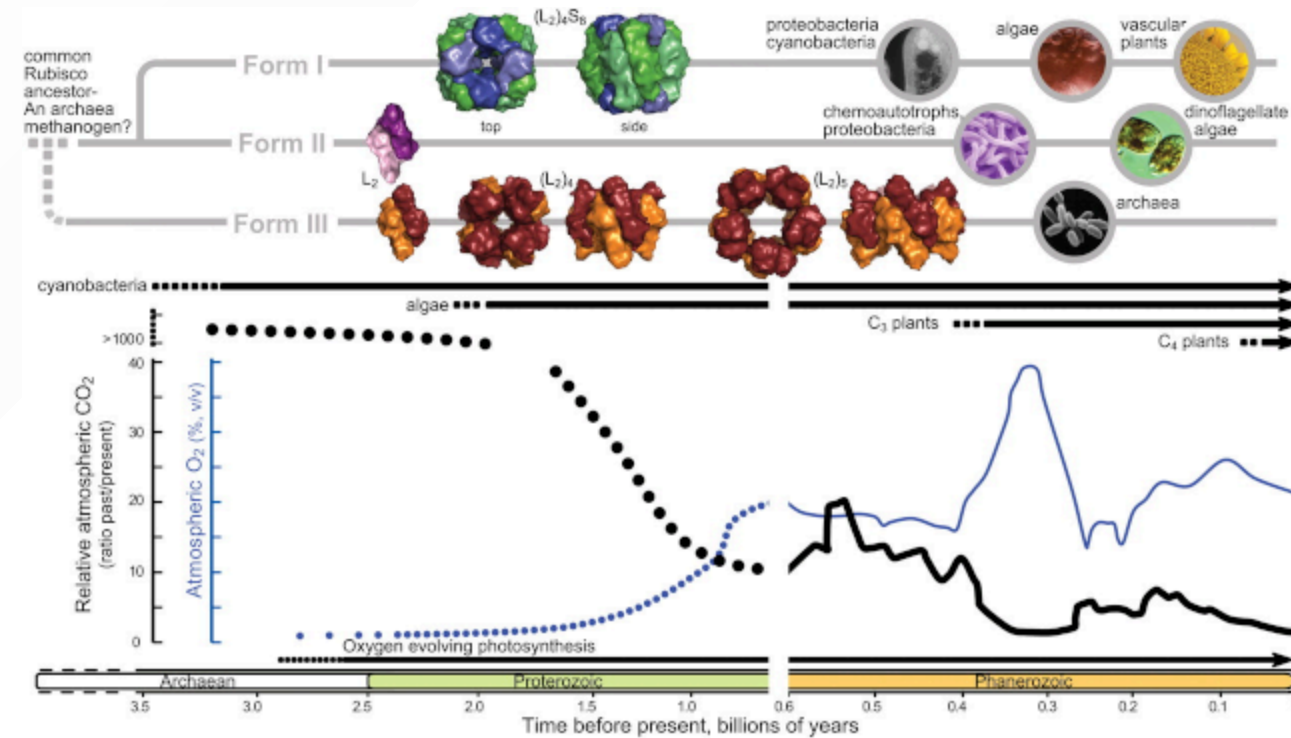
- The most abundant protein on Earth
- essential component of photosynthesis
- primary role is to convert  $\text{CO}_2$  to organic carbon



# Evolution of RuBisCO

- Sequence alignment can infer the evolutionary history of RuBisCO
- when novel features appeared
- And compare that with a geological understanding of the atmosphere at that time

ie: We can associate **features** in the protein with the **environment** in which it evolved?



# The value of understanding RuBisCO's features?

Q: Can we design more efficient RuBisCO?

↑ RuBisCO efficiency would lead to

- Crop yield ↑
- Carbon sequestration ↑

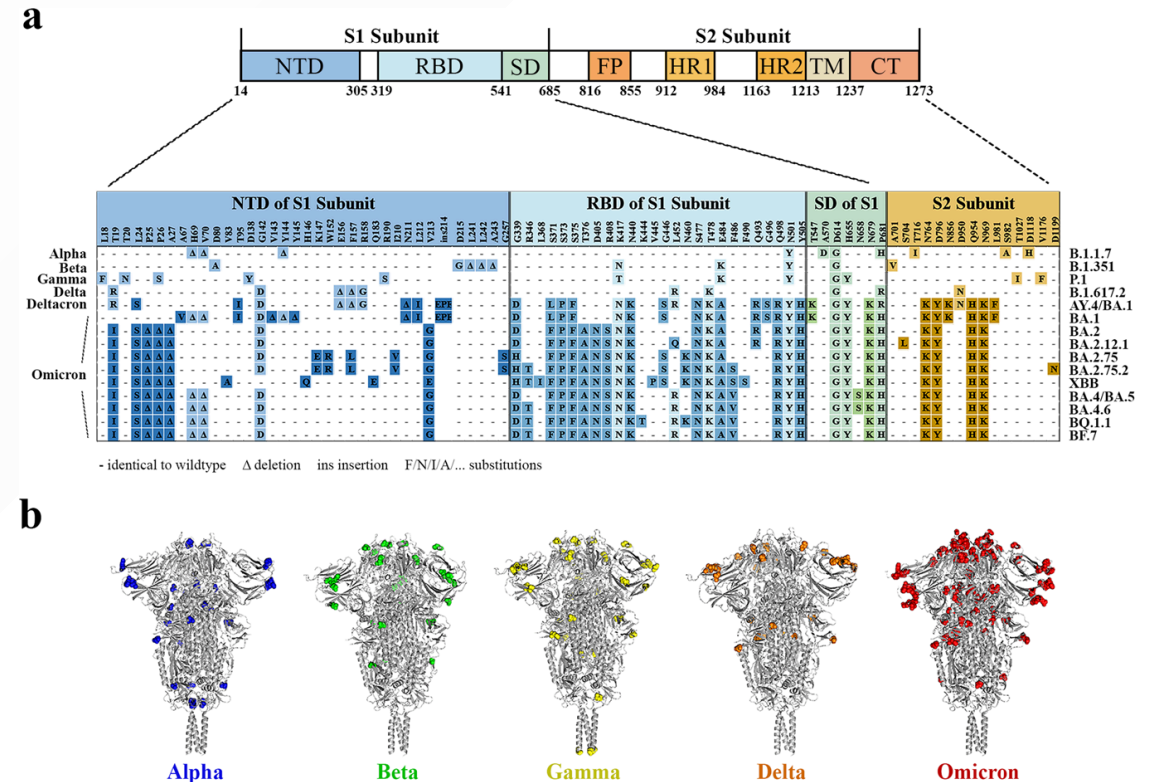
**This is a big question that sequence alignment can contribute to answering**

# Consider the spike protein of SARS-CoV-2

## Sequence alignment

- allows us to identify conserved regions for vaccine/drug development
- can help us predict the virus's trajectory

**These are big questions that sequence alignment can contribute to answering**




Alignment of S mutation points of SARS-CoV-2 variants



## Consider our immediate family

- What can we learn about our own evolution from our closest relatives?
- Can that knowledge inform biomedical science

**This is a big question that sequence alignment can contribute to answering**



The family tree of great apes

# Sequence alignment is a big job

- Historically sequence alignment was done manually, like a really big evil jigsaw puzzle
  - Since 1972 it's become a computational problem
  - The task is to compare **each** letter in **each** sequence with **all** the letters of **every** other sequence.
- 
- The terms: **each**, **all** and **every** should tell you that it will be a big job for computers too.

# Exhaustive sequence alignment takes time

A computational scientist might say that the asymptotic complexity of an exhaustive alignment is given by the big-O notation

$$O(L_1 \times L_2 \times \dots L_n) \text{ or } O(L^n)$$

Where:

- $L$  is the average length of the sequence
- $n$  is the number of sequences



“Big-O tells you how code **slows** as data **grows**” *Ned Batchelder*

# Too much math?

Let's rephrase this big-O notation as the order of  $Work(L^n)$

---

So we can reframe this as *Work* **slows** as data **grows**

Sequence length	number of sequences	<i>Work</i> required (comparisons)
1,000	2	1 Million
1,000	3	1 Billion
1,000	4	1 Trillion
1,000	5	1 Quadrillion

# Too much math?

Let's rephrase this big-O notation as the order of  $Work(L^n)$

---

So we can reframe this as *Work* **slows** as data **grows**

Sequence length	number of sequences	<i>Work</i> required (comparisons)
1,000	3	1 Billion
2,000	3	8 Billion
3,000	3	27 Billion
4,000	3	64 Billion

## The scale of our 3 big problems

Genomes	Length (bp)	Number	Work required
RuBisCO producers	1.5-500 mbp	350K <sup>1</sup>	millions <sup>hundreds of thousands</sup>
SARS-CoV-2	~29 kbp	>5M <sup>2</sup>	29 thousand <sup>5 million</sup>
Great apes	~30mbp	5	30 million <sup>5</sup>

<sup>1</sup> back of the napkin math = 300K species of **plants** + 10's of thousands of species of **algae** + thousands of species of **cyanobacter**

<sup>2</sup> 5.1M in GISAID and other public repositories as of Oct 2021 - [www.nature.com/articles/s41588-022-01033-y](https://www.nature.com/articles/s41588-022-01033-y)


**WE'RE GONNA NEED**



**MORE HAMSTERS**

# Progressive alignment

Progressive alignment is a strategy that reduces the work required






# Progressive alignment

Progressive alignment is a strategy that reduces the work required

## Strategy:

- align the 2 most closely related sequences into a statistical model called a profile




# Progressive alignment

Progressive alignment is a strategy that reduces the work required

## Strategy:

- align the 2 most closely related sequences into a statistical model called a profile
- align that profile with the next most closely related sequence




# Progressive alignment

Progressive alignment is a strategy that reduces the work required

## Strategy:

- align the 2 most closely related sequences into a statistical model called a profile
- align that profile with the next most closely related sequence
- Do that  $\binom{n}{2}$  times



# Progressive alignment

Progressive alignment is a strategy that reduces the work required

## Strategy:

- align the 2 most closely related sequences into a statistical model called a profile
- align that profile with the next most closely related sequence
- Do that  $\binom{n}{2}$  times

This reduces the order of  $Work(L^n) \rightarrow Work(n^2.L^2)$   
... which is a lot less *Work*

---

→ Profile [[[AB]C]D]

## **Progressive multiple sequence alignment (MSA)**

- **To align multiple sequences first reconstruct a phylogeny to order by distance**
- **To reconstruct a phylogeny first align all sequences**

## **Progressive multiple sequence alignment (MSA)**

- **To align multiple sequences first reconstruct a phylogeny to order by distance**
- **To reconstruct a phylogeny first align all sequences**

**Do you see the problem?**

**MSA IS USED TO CREATE PHYLOGENETIC TREES**



**PHYLOGENETIC TREES ARE USED TO GUIDE MSA**

imgflip.com

# The problem space

Sequence alignment is sensitive to


- The **length** of sequences to be aligned
  - The **number** of sequences to be aligned
  - the “ Chicken and Egg ” problem
- 

An ideal strategy would reduce

- The **length** of sequences to be aligned
- The **number** of sequences to be aligned
- Dependence on knowing the phylogeny in advance



## What if we could **quickly** remove similar regions?

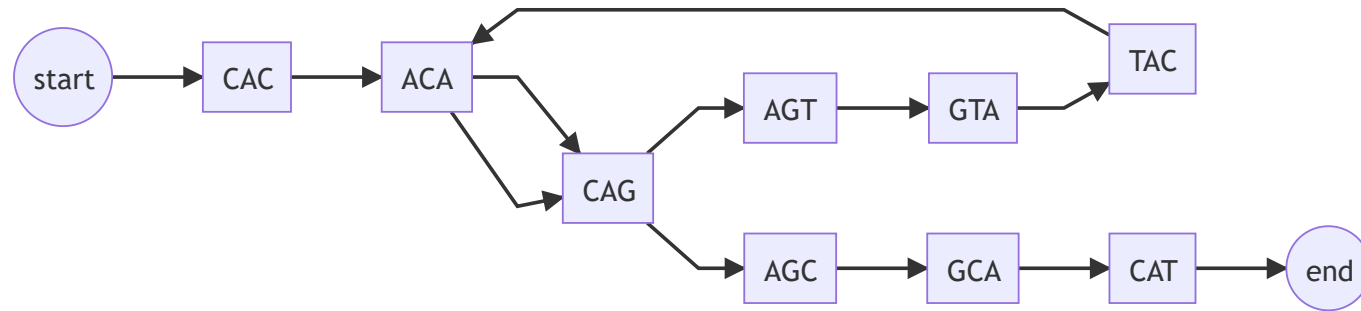


**We'd could focus resources on just the regions that differ**

# Sequence alignment using De Bruijn Graphs

My work builds upon the work by Xingjian Leng in 2022, under the supervision of Dr. Yu Lin and Prof. Gavin Huttley.

Xingjian tackled the length problem using de Bruijn graphs



You might have heard of de Bruijn graphs for sequence assembly from reads.

De Bruijn graphs can also be used for sequence alignment.

# De Bruijn graphs

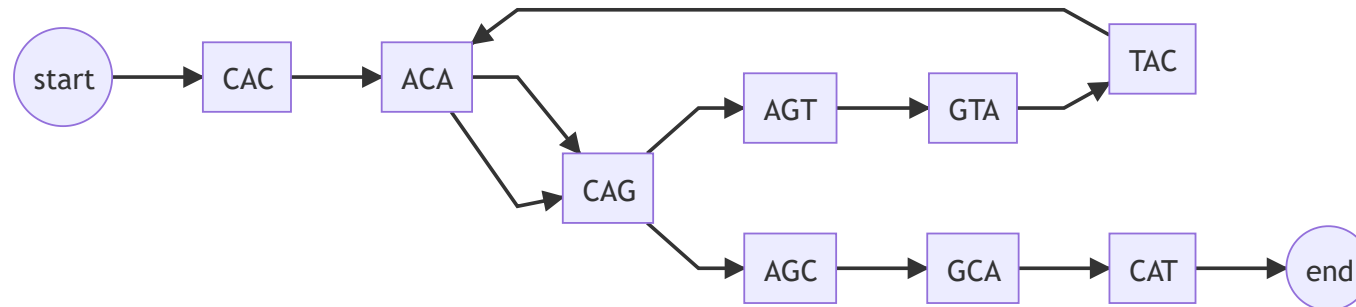
A De Bruijn graph is a directed graph that represents unique overlapping subsequences

Building a De Bruijn graph is  $Work(nL)$

**This “Work” scales linearly not exponentially.**

Consider the following sequence as a de Bruijn graph of order 3 (nodes overlap by 2 characters):

C	A	C	A	G	T	A	C	A	G	C	A	T
---	---	---	---	---	---	---	---	---	---	---	---	---



# De Bruijn graphs

A De Bruijn graph is a directed graph that represents unique overlapping subsequences

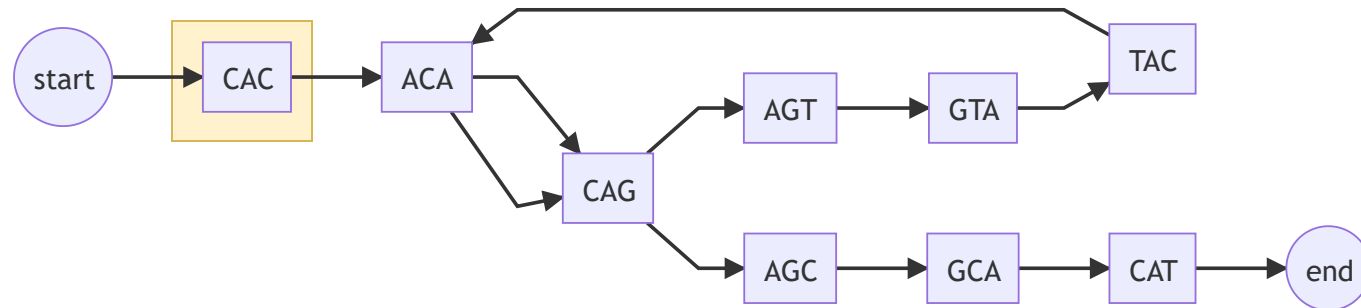
Building a De Bruijn graph is  $Work(nL)$

**This “Work” scales linearly not exponentially.**

Consider the following sequence as a de Bruijn graph of order 3 (nodes overlap by 2 characters):

C	A	C	A	G	T	A	C	A	G	C	A	T
---	---	---	---	---	---	---	---	---	---	---	---	---

C	A	C
---	---	---



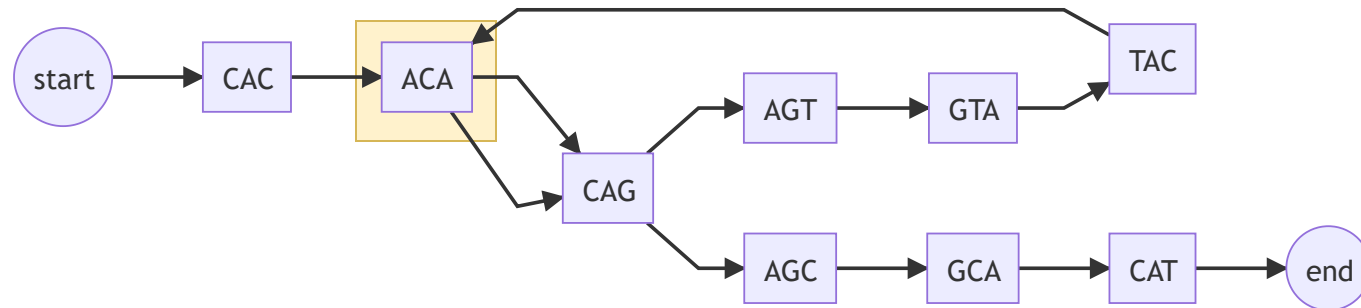

# De Bruijn graphs

A De Bruijn graph is a directed graph that represents unique overlapping subsequences

Building a De Bruijn graph is  $Work(nL)$

**This “Work” scales linearly not exponentially.**

Consider the following sequence as a de Bruijn graph of order 3 (nodes overlap by 2 characters):



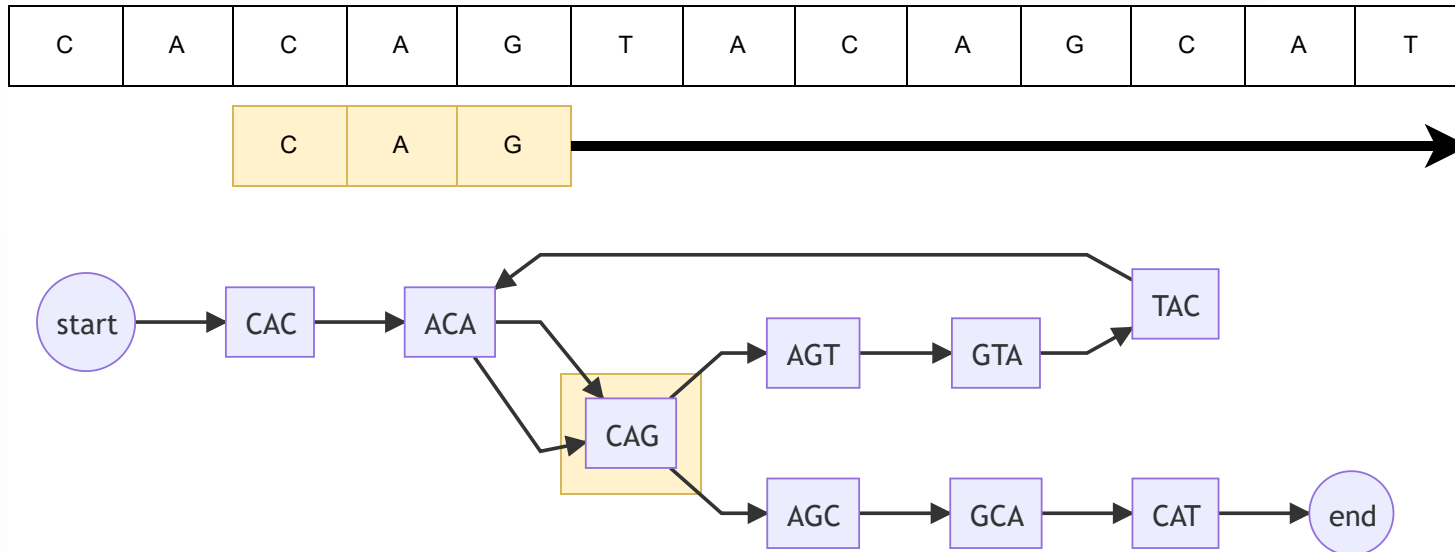
# De Bruijn graphs

A De Bruijn graph is a directed graph that represents unique overlapping subsequences

Building a De Bruijn graph is  $Work(nL)$

**This “Work” scales linearly not exponentially.**

Consider the following sequence as a de Bruijn graph of order 3 (nodes overlap by 2 characters):




And we keep going until we have created the graph

## Reducing the **length** of sequence to be aligned

Sequence A: CACAGTACGGCAT



Sequence B: CACAGTACTGCAT



These differ in just one nucleotide in the middle of the sequence


If we combine both sequences into a single de Bruijn graph, it will develop “**bubbles**” where regions are different.

# Reducing the **length** of sequence to be aligned

Sequence A: CACAGTACGGCAT




Sequence B: CACAGTACTGCAT



These differ in just one nucleotide in the middle of the sequence


If we combine both sequences into a single de Bruijn graph, it will develop “**bubbles**” where regions are different.





## Reducing the **length** of sequence to be aligned


If we transform the graph into a partial order graph, we can see the regions that are similar which we don't need to align, and the regions that are different (in the gold box) which we do.



Now we can use a traditional algorithm to align the regions  $\boxed{G}$  and  $\boxed{T}$ , and we've reduced  $Work(14 \times 14)$  down to  $Work(1 \times 1) = \mathbf{196x}$  less “work”.

# De Bruijn multiple sequence alignment

And we can extend this trivially to multiple sequences. Consider aligning 4 sequences



Now we've reduced  $\text{Work}(13 \times 13 \times 16 \times 16)$  down to  $\text{Work}(1 \times 1 \times 5 \times 4)$

= **2,163x** less “work”

## Taking the de Bruijn graph to the next level

- recall exact alignment has a Work order of  $Work(L^n)$  or  $Work(L_1 \times L_2 \times \cdots \times L_n)$
- if we reduce the length of the sequences we need to align then we reduce L

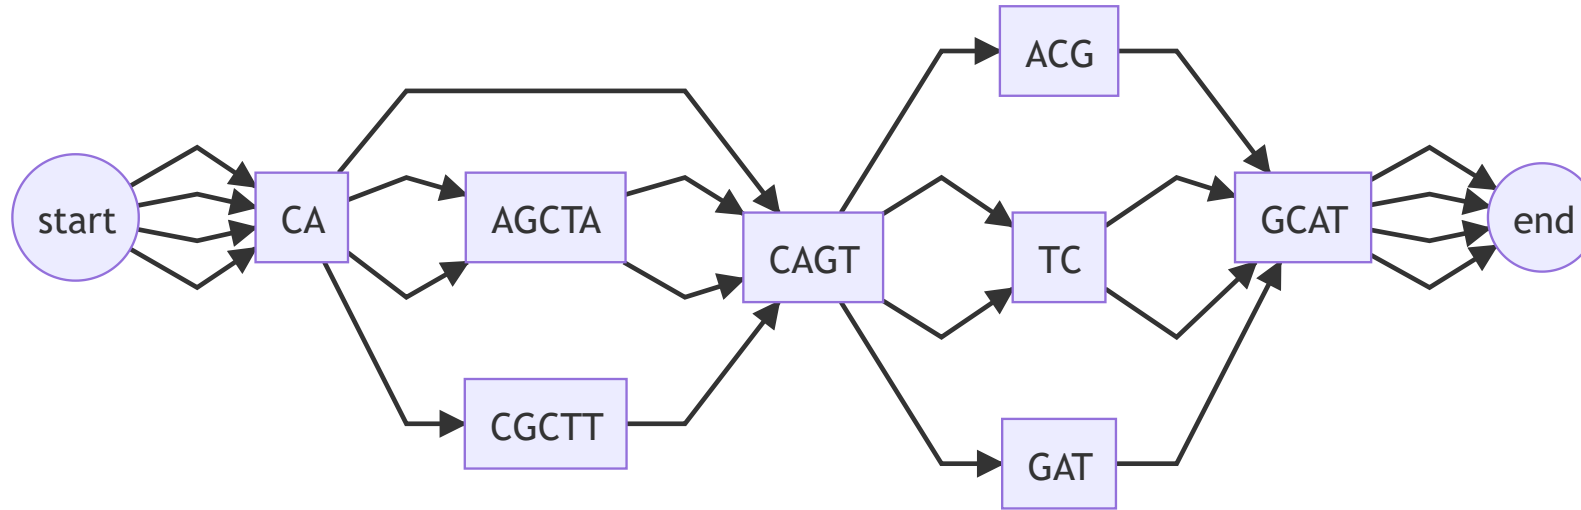
## Taking the de Bruijn graph to the next level

- recall exact alignment has a Work order of  $Work(L^n)$  or  $Work(L_1 \times L_2 \times \cdots \times L_n)$
- if we reduce the length of the sequences we need to align then we reduce L

Can we change the **number** of sequences to align?

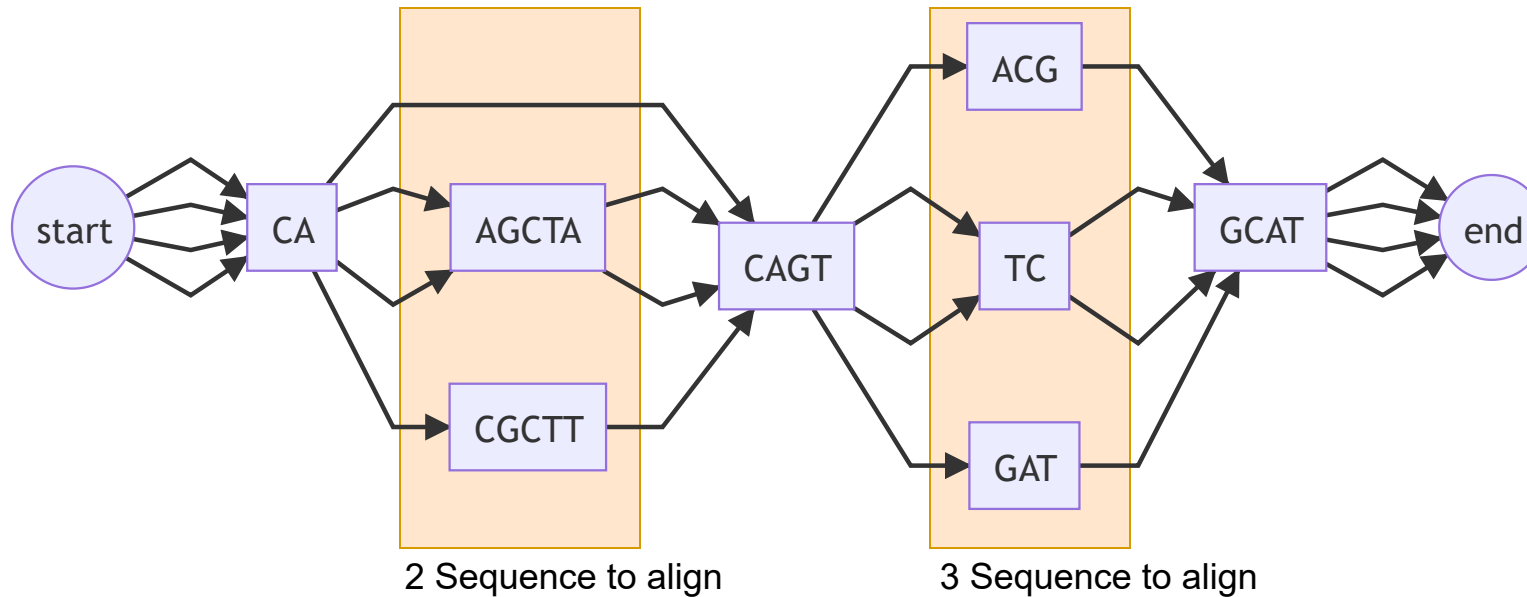
## Reducing the **number** of sequences to be aligned

Consider this partial order graph containing 4 sequences with overlaps removed



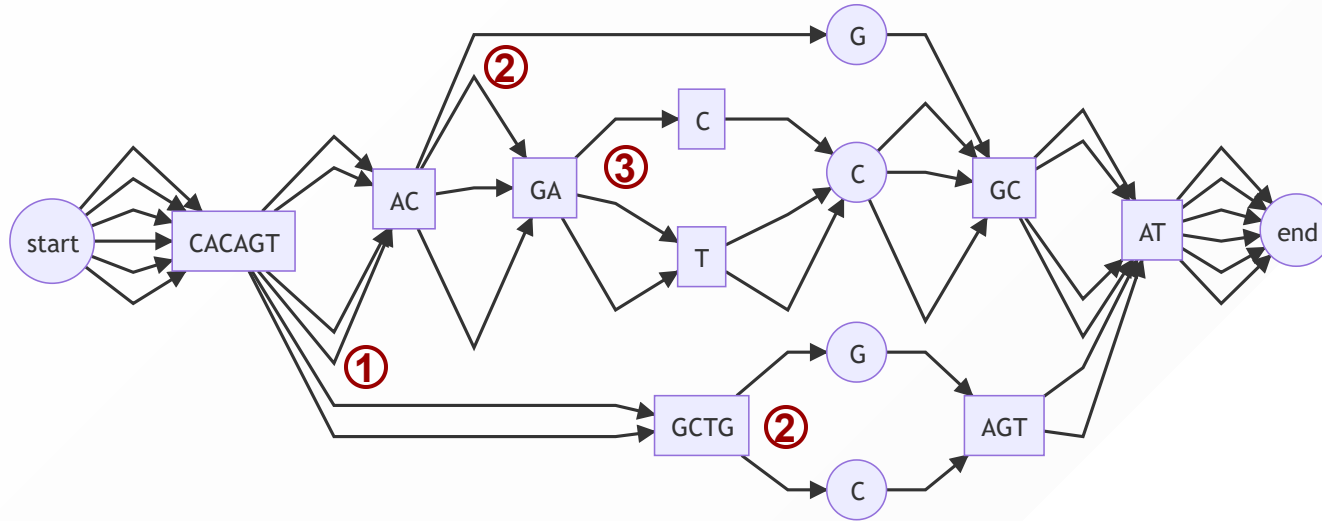
# Reducing the **number** of sequences to be aligned

Consider this partial order graph containing 4 sequences




exhaustive alignment	reduce length	reduce length & number
$Work(13 \times 14 \times 17 \times 17)$	$Work(5^4 + 3 \times 2 \times 2 \times 3)$	$Work(5^2 + 3 \times 2 \times 3)$
52,598	661 ( <b>79x</b> vs exhaustive)	43 ( <b>1223x</b> vs exhaustive)

## Reduce the dependence on the phylogeny





By ordering progressive alignment by descending “bubble” depth, we can progressively align without needing to know in advance the phylogenetic relation between sequences.

## Reduce the dependence on the phylogeny

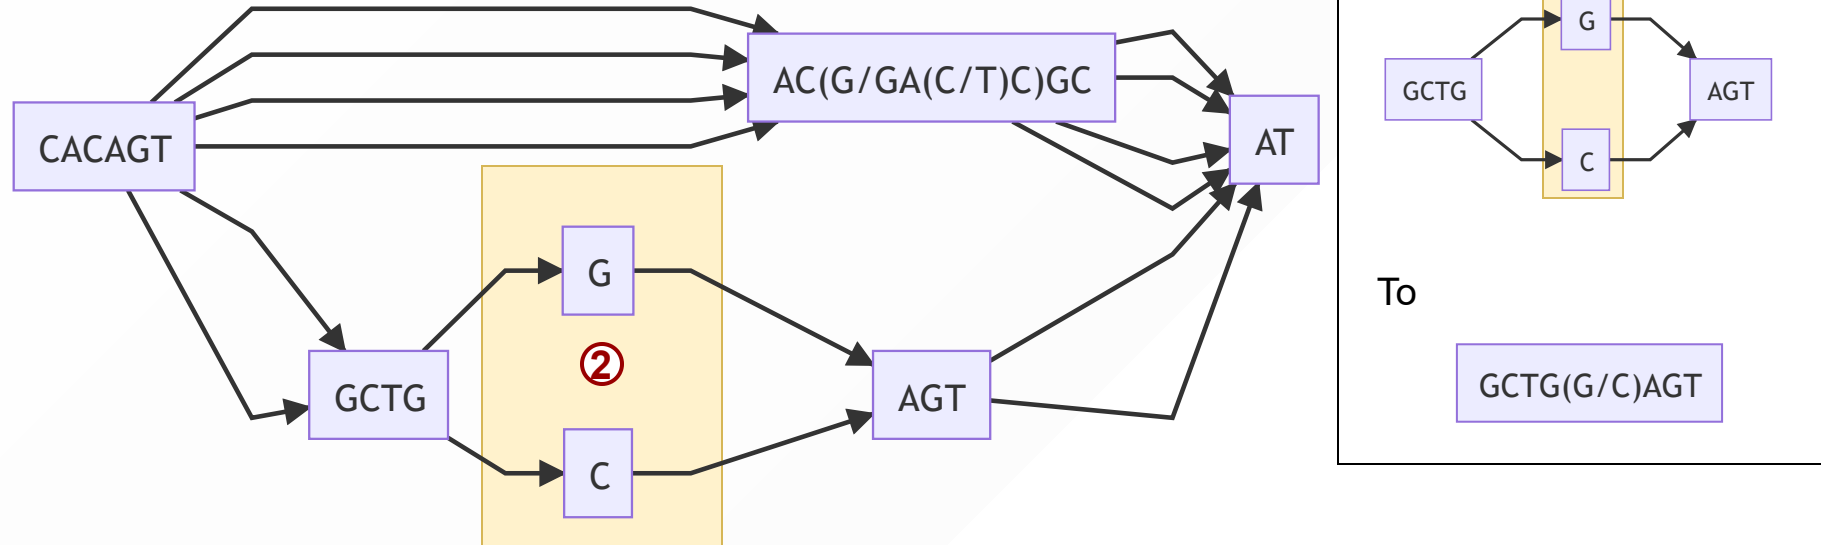




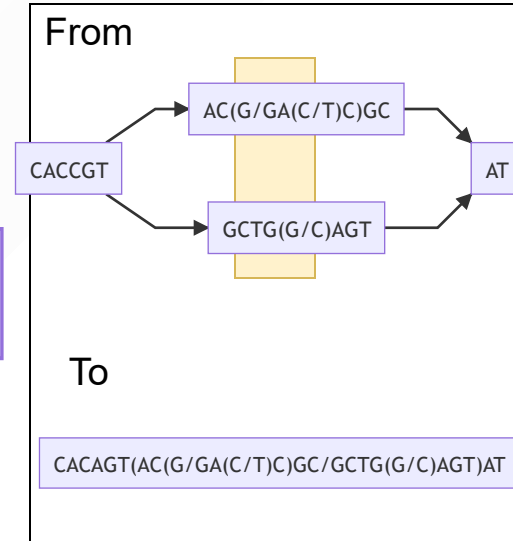
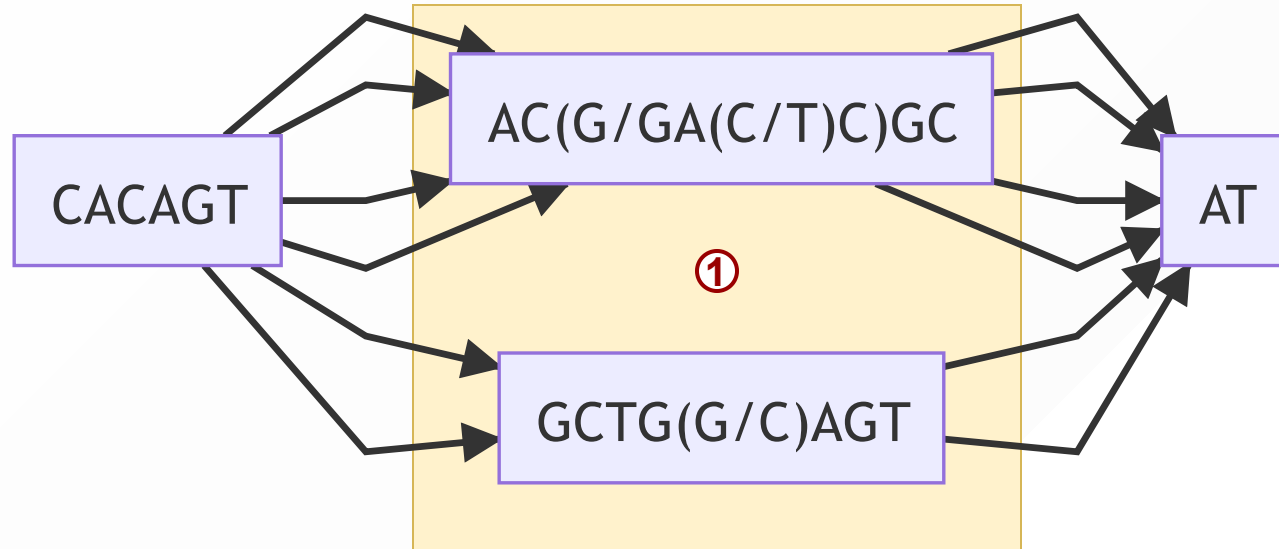
## Reduce the dependence on the phylogeny



# Reduce the dependence on the phylogeny



## Reduce the dependence on the phylogeny



## Reduce the dependence on the phylogeny

CACAGT(AC(G/GA(C/T)C)GC/GCTG(G/C)AGT)AT


**Alignment completed without requiring a phylogenetic guide tree**

## Project aims

- Investigate De Bruijn graphs for multi-sequence alignment (MSA)
- Build a python library to
  - Resolve the De Bruijn graph to a partial order graph of segments to align
  - identify “bubbles”
  - Develop unit tests to verify correctness
- Develop statistics for de Bruijn graphs to predict efficiency

## Results: Work statistic

Consider the same order 3 de Bruijn graph




- Work calculates the order of alignment work using 4 strategies
  - Exact =  $13 \times 12 \times 9 = 1404$
  - Progressive =  $13 \times 12 + 13 \times 9 = 285$
  - DBG\_L =  $4 \times 5 + 5 \times 1 = 25$  (simplification of sequence length)
  - DBG\_LN =  $0 \times 1 + 5 \times 1 = 5$  (simplification of sequence length and count)

## Results: Quickwork statistic

Consider this order 3 de Bruijn graph containing 3 sequences [ **ACAGTACGGCAT**, **ACAGTACTGGCAT**, **ACAGCGCAT** ] of length 12, 13 and 9

When transformed into a partial order graph



Contains the following nodes (left to right) with ACAG and GCAT (already aligned sections) removed TAC+T+G+C

Quickwork =  $\sum$  non aligned node length = 6

Quickwork has an order of  $Work(node\_count)$

# Results: Calculated order complexity

## Sample Sequence sets

Genomes	Average Length (bp)	Number	Description
BRCA1_divergent	~3k	56	BRCA1 gene divergent sample of 7 chosen from among 56 mammal species
BRCA1_hominae	~3k	4	BRCA1 gene from 4 hominae
SARS-CoV-2	30k	22	SARS-CoV-2 genomes (citation needed)
IBD_phage	40k	60	IBD phage components ( <a href="https://doi.org/10.1016/j.cell.2015.01.002">https://doi.org/10.1016/j.cell.2015.01.002</a> )
Ocean_phage	40k	130	Tara oceans phage components ( <a href="https://doi.org/10.1126/science.1261605">https://doi.org/10.1126/science.1261605</a> )

Work metric	Description
Exact	Exhaustive alignment
Progressive	Progressive alignment
DBG(3)_L	De Bruijn graph of order 3 with length simplification
DBG(3)_LN	De Bruijn graph of order 3 with length and number simplification
DBG( $k$ )_LN	De Bruijn graph of order $k$ with length and number simplification



## Results: Calculated order of Work

Genomes	Exact	Progressive	DBG(3)_L	DBG(3)_LN	DBG(4)_LN	DBG(5)_LN	DBG(6)_LN	DBG(7)_LN	DBG(8)_LN
BRCA1_divergent									
BRCA1_hominae									
SARS-CoV-2									
IBD_phage									
Ocean_phage									

## Results: Quickwork

Genomes	dBG(3)	dBG(4)	dBG(5)	dBG(6)	dBG(7)	dBG(8)	dBG(9)
BRCA1_divergent							
BRCA1_hominae							
SARS-CoV-2							
IBD_phage							
Ocean_phage							

## Sample unit tests: cyclic sequences

```
def test_pog_cycle(output_dir: Path):
    dbg = dbg_align.DeBruijnGraph(3, cogent3.DNA)
    dbg.add_sequence({
        "seq1": "ACAGTACGGCAT",
        "seq2": "ACAGTACTGGCAT",
        "seq3": "ACAGCGCGCAT" # contains cycle
    })
    with open(output_dir / "cycle.md", "w") as f:
        f.write("```mermaid\n")
        f.write(dbg.to_mermaid())
        f.write("```")
    assert dbg.has_cycles()
    assert len(dbg) == 3
    assert dbg.names() == ["seq1", "seq2", "seq3"]
    assert dbg["seq3"] == "ACAGCGCGCAT" # contains cycle

    dbg.to_pog()
    # write mermaid out to testout folder
    with open(output_dir / "cycle_compressed.md", "w") as f:
        f.write("```mermaid\n")
        f.write(dbg.to_mermaid())
        f.write("```")
```

# Discussion

de Bruijn graphs offer an interesting method to

- Reduce the impact of both sequence length and sequence number over traditional alignment approaches
- Break the tautology at the heart of both Sequence alignment, and Phylogenetic reconstruction

**This method may make some very big questions tractable**

# Future directions

Investigate the potential of using de Bruijn Graphs to;

- Identify reverse complimented regions from a dBG
- Identify genetic distance and infer phylogeny from a dBG
- Process sequences in databases storing dBG structures back to the database, reducing active memory limits for large numbers of large sequences
- Investigate advantage wrt species subject to lateral gene flow
  - eg: Bacteria, Archaea
  - identifying multi-rooted phylogenies
- Investigate using dBG's for targeted sequence extraction using pattern recognition templates
- Identify strategies for choosing the ideal kmer length for a given set of sequences

# Thanks

- Gavin Huttley
- Yu Lin
- Vijini Mallawaarachchi
- Xinjian Leng

## ... and the Huttleylab



# Thanks

- Gavin Huttley
- Yu Lin
- Vijini Mallawaarachchi
- Xinjian Leng

## Q&A

## ... and the Huttleylab



# Errata

“ Abandon all hope ye who pass this point

*Tolkein ... probably*

”

- [Sequence alignment order complexity](#)
- [Pairwise sequence alignment methods](#)
- [Multiple sequence alignment strategies](#)
- [unit tests](#)



# Sequence alignment order complexity

[<<Back to Errata](#)

## Pairwise sequence alignment

- Compare every letter in one sequence to every letter in the other
- order complexity of  $O(mn)$ 
  - where **m** and **n** are lengths of the sequences

## Multiple sequence alignment (MSA)

- Perform a pairwise alignment of every sequence to every other sequence
- order complexity of  $O(L^n)$ 
  - where **L** is the length of the sequences
  - **n** is the number of sequences

# Pairwise sequence alignment methods: $O(mn)$

[<<Back to Errata](#)

- Needleman-Wunsch algorithm: global alignment for highly similar sequences
  - scoring system that penalises gaps and mismatches
- Smith-Waterman algorithm: better for local alignment to find conserved domains
  - allows for alignment to reset when the score falls to 0

Compare each nucleotide in one sequence to each nucleotide in the other sequence

Given a simple scoring system +1 match, -1 mismatch, -2 gap ( $\delta$ )

Where  $F(i, j) = \max$  of the following

$$\begin{aligned} & \nwarrow F(i-1, j-1) + s(A_i, B_j), \quad (\text{match/mismatch}) \\ & \uparrow F(i-1, j) + \delta, \quad (\text{deletion}) \\ & \leftarrow F(i, j-1) + \delta, \quad (\text{insertion}) \end{aligned}$$

	gap	A	G	C	A	A
gap	0	$\leftarrow -2$	$\leftarrow -4$	$\leftarrow -6$	$\leftarrow -8$	$\leftarrow -10$
A	$\uparrow -2$	$\nwarrow \mathbf{1}$	$\leftarrow -1$	$\leftarrow -3$	$\leftarrow -5$	$\nwarrow \nwarrow -7$
C	$\uparrow -4$	$\uparrow -1$	$\nwarrow 0$	$\nwarrow \mathbf{0}$	$\leftarrow -2$	$\leftarrow -4$
G	$\uparrow -6$	$\uparrow -3$	$\nwarrow 0$	$\nwarrow -1$	$\nwarrow \mathbf{1}$	$\leftarrow -1$
A	$\uparrow -8$	$\nwarrow \uparrow -5$	$\uparrow -2$	$\nwarrow -1$	$\uparrow -1$	$\nwarrow \mathbf{2}$
A	$\uparrow -10$	$\nwarrow \uparrow -7$	$\uparrow -4$	$\nwarrow \uparrow -3$	$\nwarrow -2$	$\nwarrow \uparrow \mathbf{0}$

# Multiple sequence alignment (MAS) strategies

[<<Back to Errata](#)

- Pairwise alignment of each possible pair
  - $\binom{n}{2} \times O(L^2) = \frac{n(n-1)}{2} \times O(L^2) = O(n^2.L^2)$
- Progressive alignment eg: ClustalW
  - create a guide tree
  - Progressively align pairs most closely related to profiles, and then align profiles
- Iterative methods eg: MUSCLE, T-Coffee, MAAFT
  - create an preliminary fast less accurate alignment
  - iteratively improve alignment using some scoring function
  - Complete when some convergence criterion is met
- Hidden markov models  $O(nL) + O(LM)$  (M is the number of states in the model)
  - eg: HMMER
  - create a statical model of the transition between states
  - Determine likely alignment based on the model

# Unit tests

[<<Back to Errata](#)

library against edge case sequence alignments

- \* long sequences
- \* numerous sequences
- \* cyclic sequences
- \* bubbles within bubbles
- \* sequential bubbles